

Neural Prophet

Extendable and Scalable Forecasting

Oskar Triebel, Stanford University
Facebook Forecasting Summit, October 5th 2020

Oskar Triebe

Raised in Switzerland

BS, ETH Zurich

MS, PhD, Stanford

Vectra AI

Merantix

Luminovo

The outdoors

DIY power electronics

Digital human rights



Motivation

What is (Neural)Prophet?

Ar-Net Module

Development Story

Hands-On: NeuralProphet

Outlook

Forecasting Landscape

Traditional Methods

(S)ARIMA(X)

(V)ARMA(X)

GARCH

(S)Naïve

Gaussian Process

HMM

(T)BATS

Seasonal + Trend Decomposition

Exponential Smoothing

Holt-Winters

Dynamic Linear Models

Prophet

Deep Learning

AR-Net

ES-RNN

LSTM

N-BEATS


Transformer

WaveNet

CNN

Other ML

Chasm



Bridge

- Beginner-friendly
- Interpretable
- Scalable
- Customizable Model
- Extendable Code
- Workflow Integration

Primer on Facebook Prophet




Forecasting at Scale

Sean J. Taylor*†
Facebook, Menlo Park, California, United States

sjt@fb.com

and


Benjamin Letham†
Facebook, Menlo Park, California, United States
bletham@fb.com



<https://peerj.com/preprints/3190/>

Prophet models interpretable time series components.

(Neural)Propet



$$y(t) = g(t) + s(t) + h(t) + \epsilon_t.$$

$g(t)$ Trend

$h(t)$ Events

$s(t)$ Seasonality





$x(t)$ Regressors

$g(t)$

$h(t)$

$s(t)$

$x(t)$





Stan



PyTorch

Extensible code.

Models suited for larger datasets.

Lagged input variables.

... and adds:

- Auto-Regression
- Covariates (lagged regressors)
- Configurable nonlinear deep layers
- Tuneable to specific forecast horizons
- Custom metrics and losses
- Workflow tools




Extensibility:

- Anything trainable by gradient descent can be added as module

How did we get here?

Let's start at the beginning...

TS-Cell



simple
interpretable
performant


Vision at outset:
Build an easy to use general time-series cell.

Focusing on Auto-Regression

$$y_t = c + \sum_{i=1}^{i=p} w_i * y_{t-i} + e_t \quad \longrightarrow \quad \min_{\theta} L(y, \hat{y}, \theta)$$


Simplest form: AR-Net(0)


Interpretable




General purpose forecaster: AR-Net(n)

Stronger modeling ability






Closeness to true coefficients



$$sTPE = 100 \cdot \frac{\sum_{i=1}^{i=p} |\hat{w}_i - w_i|}{\sum_{i=1}^{i=p} |\hat{w}_i| + |w_i|}$$


Forecast accuracy




$$MSE = \frac{1}{n} \sum_1^n (y_t - \hat{y}_t)^2$$

AR-Net is quadratically faster for large p values.

AR-Net



Sparse AR process



Trained with SGD (Adam)


$$\min_{\theta} \quad L(y, \hat{y}, \theta) + \lambda(s) \cdot R(\theta)$$
$$\lambda(s) = c_\lambda \cdot (s^{-1} - 1)$$

$$s = \frac{\hat{p}_{data}}{p_{model}}$$
$$c_\lambda \approx \frac{\sqrt{\hat{L}}}{100}$$

$$R(\theta) = \frac{1}{p} \sum_{i=1}^p \frac{2}{1 + \exp(-c_1 \cdot |\theta_i|^{\frac{1}{c_2}})} - 1$$

Yes, we can automatically learn sparse AR-Coefficients.


AR-Net




Sparse AR-Net surpasses Classic AR and scales to large orders.

AR-Net

Closeness to true coefficients




MSE loss on forecast target




```
any time series
learn = init_ar_learner(
    series=df,
    ar_order=3,
    n_forecasts=1,
    valid_p=0.25,
    sparsity=1.0/2.0,
)
```

```
lr = learn.lr_find()[0]
learn.fit_one_cycle(n_epoch=10, lr)
print(coeff_from_model(learn.model))
```



epoch	train_loss	valid_loss	mae	time
0	0.091714	0.037677	0.165344	00:00
1	0.043905	0.019270	0.112051	00:00
2	0.026220	0.010630	0.083301	00:00
3	0.018274	0.011065	0.084113	00:00
4	0.014546	0.010218	0.081142	00:01
5	0.013065	0.009898	0.080138	00:00
6	0.012065	0.010537	0.083343	00:01
7	0.011674	0.011297	0.085578	00:00
8	0.011395	0.010602	0.082547	00:00
9	0.011284	0.011390	0.085972	00:00
10	0.011243	0.011658	0.086935	00:00



building on open-source Fastai2, PyTorch, Python

```
ar coeff [[0.19132832, 0.3154733, -0.5822663, -0.049190696, 0.0042145597, 0.0015144324]]
```

Paper on Arxiv

<https://arxiv.org/abs/1911.12436>

The screenshot shows the Arxiv paper page for "AR-Net: A simple Auto-Regressive Neural Network for time-series". The page includes the title, authors (Oskar Triebe, Nikolay Laptev, Ram Rajagopal), submission date (27 Nov 2019), abstract, and several sections of the paper's content. On the right side, there is a sidebar with download options (PDF, Other formats), current browse context (cs.LG), change to browse by (cs, stat, stat.ML), references & citations, export citation, bookmark, and BibTeX links.

Code on Github

<https://github.com/ourownstory/AR-Net>

The screenshot shows the Github repository page for "ourownstory / AR-Net". It displays the repository's overview, including the master branch (1 branch, 15 commits), recent activity, and a list of commits. The commits are:

Commit	Message	Time
712bbe2	ourownstory update Readme	on Jun 11
v0_1	releasing v1.0 of AR-Net with simpler use, based on fastai2	7 months ago
v1_0	updated notebooks with cells run	6 months ago
AR-Net_paper.pdf	add paper to files	8 months ago
README.md	update Readme	4 months ago


Annotations on the screenshot highlight specific commits:

- A curved arrow points from the "v0_1" commit to the text "v 0.1: plain PyTorch".
- A curved arrow points from the "v1_0" commit to the text "v 1.0 built on Fastai2".

Developing an open-source package


For the first time.

Thank you to all open-source contributors, you have left us a huge legacy!




Interview lots of people. Prototype. Iterate.

Dev Story



Welcome to fastai v2

NB: This is still in early development. Use v1 unless you want to contribute to the next version of fastai



*That's actually what
I ended up doing ...*



... and back to plain PyTorch

Design challenges.

Dev Story

Pun intended.



User



Coder

Explicit




Implicit

bug

- \(°_o) / -

A car is more than an engine.

New Vision



Preprocessing



Diagnostics



Deployment

That's why we will include helpful tools to streamline your workflow.

Beta NeuralProphet


Model components

Code snippets & plots

S	<i>Seasonality</i>	
T	<i>Trend</i>	
E/H	<i>Events / Holidays</i>	
?	<i>Uncertainty</i>	
-	<i>Sparsity / Regularization</i>	Priors → Regularization
AR	<i>Autoregression</i>	AR-Net
X	<i>Covariates / Regressors</i>	Lagged and future-known
NN	<i>Nonlinear (deep) layers</i>	

Piecewise linear trend

- N changepoints
- Segmentwise independent
- Automatic changepoint detection
- Optional logistic growth

$$g(t) = (k + \mathbf{a}(t)^\top \boldsymbol{\delta})t + (m + \mathbf{a}(t)^\top \boldsymbol{\gamma}),$$


Seasonality

- N Fourier terms
- Automatic yearly, weekly, daily

$$s(t) = \sum_{n=1}^N \left(a_n \cos\left(\frac{2\pi n t}{P}\right) + b_n \sin\left(\frac{2\pi n t}{P}\right) \right)$$

$$s(t) = X(t)\boldsymbol{\beta}.$$

$$X(t) = \left[\cos\left(\frac{2\pi(1)t}{365.25}\right), \dots, \sin\left(\frac{2\pi(10)t}{365.25}\right) \right]$$

Events / Holidays

- Automatic for given country
- Various user-defined formats

$$h(t) = Z(t)\kappa.$$

$$Z(t) = [\mathbf{1}(t \in D_1), \dots, \mathbf{1}(t \in D_L)]$$

Future-Known Regressors

- Single weight

Lagged Covariates

- AR-Net (y as target)

Auto-Regression

- AR-Net

Beta NeuralProphet

Model components


Code snippets & plots

Few lines from data to forecast.

NeuralProphet

```
# linear time-dependent model
m = NeuralProphet()
metrics = m.fit(df)
future = m.make_future_dataframe(df, future_periods=365)
forecast = m.predict(future)
fig_fcst = m.plot(forecast[-730:])
```

Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.



```
import pandas as pd
from neuralprophet.neural_prophet import NeuralProphet

df = pd.read_csv('../data/example_wp_log_peyton_manning.csv')




n_forecasts=1,
n_lags=0,
n_changepoints=5,
learning_rate=1.0,
loss_func='Huber',
normalize_y=True,
num_hidden_layers=0,
d_hidden=None,
ar_sparsity=None,
trend_smoothness=0,
trend_threshold=False,
yearly_seasonality='auto',
weekly_seasonality='auto',
daily_seasonality='auto',
seasonality_mode='additive',
seasonality_reg=None,
data_freq='D',
impute_missing=True,
verbose=False,
```

```
# or evaluate while training
m = NeuralProphet()
metrics = m.fit(df, validate_each_epoch=True, valid_p=0.2)
metrics.tail()
```

Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.

	SmoothL1Loss	MAE	RegLoss	SmoothL1Loss_val	MAE_val
35	0.163102	0.371323	0.0	0.485371	0.779465
36	0.161851	0.369609	0.0	0.368921	0.648736
37	0.161122	0.369219	0.0	0.366328	0.648230
38	0.168598	0.376638	0.0	0.348269	0.627886
39	0.167961	0.375777	0.0	0.362161	0.642699

```
# split manually
m = NeuralProphet()
df_train, df_val = m.split_df(df, valid_p=0.2)
train_metrics = m.fit(df_train)
val_metrics = m.test(df_val)
```





```
m = NeuralProphet(  
    n_changepoints=100,  
    trend_smoothness=2,  
    yearly_seasonality=False,  
    weekly_seasonality=False,  
    daily_seasonality=False,  
)  
metrics = m.fit(df)
```

Seasonality

NeuralProphet

```
m = NeuralProphet(  
    yearly_seasonality=16,  
    weekly_seasonality=8,  
    daily_seasonality=False,  
    seasonality_reg=1,  
)  
metrics = m.fit(df)
```

NOTICE: A Regularization strength for the seasonal Fourier Terms was set. Please note that this feature is experimental.



Events can be added in different forms.

NeuralProphet

```
superbowls = pd.DataFrame({'event': 'superbowl',
 'ds': pd.to_datetime(['2010-02-07', '2014-02-02', '2016-02-07'])})
```

```
events_df = pd.concat((superbowls, playoffs))
```

```
m = NeuralProphet()
```

```
m = m.add_country_holidays("US")
```



```
m = m.add_events(
    ["superbowl", "playoff"],
    lower_window=-1,
    upper_window=1,
    mode="multiplicative",
    regularization=0.5
)
```

```
history_df = m.create_df_with_events(df, events_df)
metrics = m.fit(history_df)
```

Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.

```
future = m.make_future_dataframe(
    history_df,
    events_df,
    future_periods=30,
    n_historic_predictions=30
)
```

```
forecast = m.predict(future)
```



Lagged Covariates / Auto-Regression

NeuralProphet

```
m = NeuralProphet(  
    n_forecasts=30,  
    n_lags=60,  
    ar_sparsity=0.1,  
    yearly_seasonality=False,  
    weekly_seasonality=False,  
    daily_seasonality=False,  
)
```


```
df = pd.read_csv('../data/example_wp_log_peyton_manning.csv')  
df['A'] = df['y'].rolling(7, min_periods=1).mean()  
df['B'] = df['y'].rolling(60, min_periods=1).mean()  
m = m.add_covariate(name='A')  
m = m.add_regressor(name='B')  
  
metrics = m.fit(df)  
future = m.make_future_dataframe(df, n_historic_predictions=30)  
forecast = m.predict(future)
```

Future-known regressors are similar

Tune .


```
m = NeuralProphet(  
    n_forecasts=30,  
    n_lags=60,  
    learning_rate=1.0,  
    loss_func='Huber',  
    normalize_y=True,  
    num_hidden_layers=2,  
    d_hidden=64,  
)
```

```
fig_comp = m.plot_components(forecast[60:])
```



Lagged Covariates / Auto-Regression: Interpret.


NeuralProphet





Focus on a specific forecast.


NeuralProphet

```
fig_prediction = m.plot_last_forecast(forecast[60:],  
include_previous_forecasts=2)
```



```
fig_param = m.highlight_nth_step_ahead_of_each_forecast(30).  
plot_parameters()
```





https://github.com/ourownstory/neural_prophet

Continues the legacy of Prophet in PyTorch.

Extensible to future state-of-the-art in forecasting.

Workflow-aware with tools for diagnostics and more.

Outlook



Next:

Presenting at ISF

Public Beta release (code + journal paper)

Future:

V 1.0 Workflow utilities, robustify training, model extensions

V 2.+ Modularize, integrate other frameworks (data, deployment)

V 3.+ Web GUI, port to other languages

THANK YOU, dear collaborators and advisors!

Team



Nikolay Laptev



Oskar Triebe



Hansika Hewamalage



Italo Lima




Caner Komurlu



Gonzague Henri



Ram Rajagopal



Riley De Haan




Christoph Bergmeir

facebook

TOTAL

Stanford University

MONASH
University



Connect with me: linkedin.com/in/oskartriebe/
github.com/ourownstory
triebe [at] stanford.edu

Thank You