

Neural Prophet

A simple time series forecasting framework

Oskar Triebe, Mateus De Castro Ribeiro, Abishek Sriramulu, Hansika Hemawalage
Stanford University, Monash University

Nov 16th, 2021

Neural Prophet

is

an open-source forecasting library.

tl;dr

Prophet in PyTorch + AR + Covar + NN + multistep + ...

Task:

Forecasting.

Data:

1E+2 to 1E+6 of samples. Unidistant, real-valued.

Dynamics:

Future values must depend on past observations.
e.g. Seasonal, trended, events, correlated variables.

Applications:

Human behavior, energy, traffic, sales, environment, server load, ...

Motivation

Model

Example

Global Modelling

Attention

Conclusion

Time series forecasting is messy. We need hybrid models to bridge the gap.

Traditional Methods

(S)ARIMA(X)
(V)ARMA(X)

GARCH

(S)Naïve

Gaussian
Process

HMM

(T)BATS

Seasonal + Trend
Decomposition

Exponential
Smoothing

Holt-Winters

Dynamic Linear
Models

Neural Prophet

Prophet

AR-Net

ES-RNN

LSTM

Transformer

DeepAR

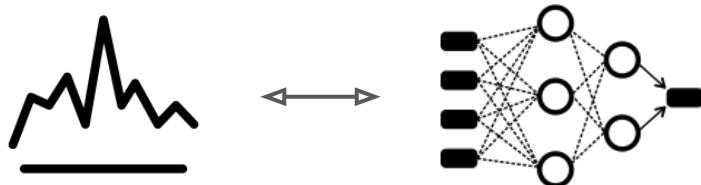
N-BEATS

WaveNet

Causal
Convolutions

Other ML

Before



Need expertise in both
time series & machine learning

NeuralProphet



Abstracts time series
& machine learning knowledge

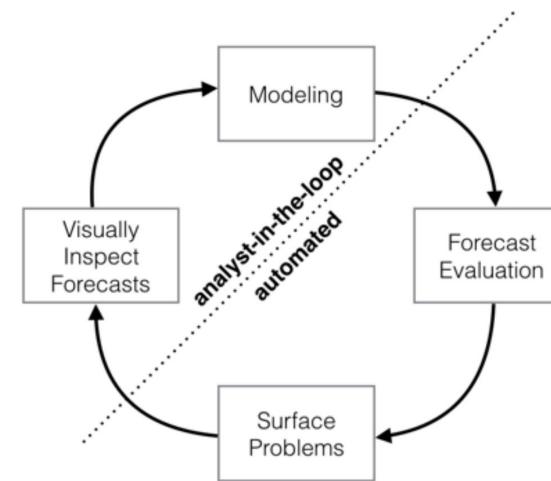
Facebook Prophet is still the most used forecasting package.



Quick from data to predictions.

Gentle learning curve.

Customizable.



Taylor, S. J., & Letham, B. (2017).
Forecasting at scale, PeerJ.
<https://peerj.com/preprints/3190/>



Prophet has three major shortcomings:

1. Missing local context for predictions
2. Acceptable forecast accuracy
3. Framework is difficult to extend (Stan)



NeuralProphet solves these:

1. Support for auto-regression and covariates.
2. Hybrid model (linear <> Neural Network)
3. Python package based on PyTorch using standard deep learning methods.

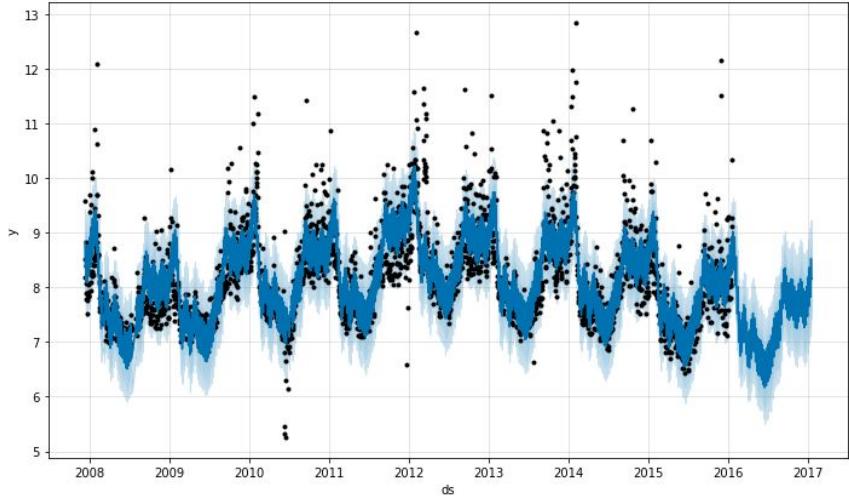


Current Model Components

S	Seasonality	-	Sparsity / Regularization
T	Trend	NN	Nonlinear (deep) layers
E/H	Events / Holidays	{ }	Global Modelling
X	Regressors	?	Uncertainty
AR	Autoregression		
Cov	Covariates		

Prophet is interpretable and decomposable. So is NeuralProphet.

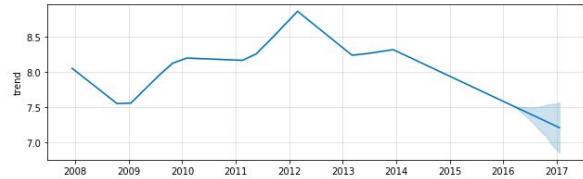
Prophet



$$y(t) = g(t) + s(t) + h(t) + \epsilon_t.$$

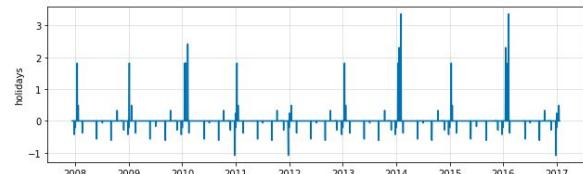
Trend

$$g(t)$$



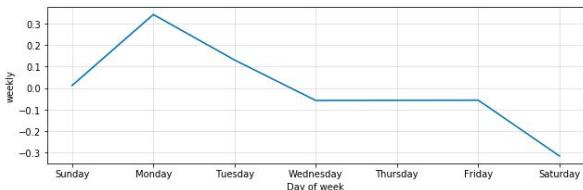
Events

$$h(t)$$



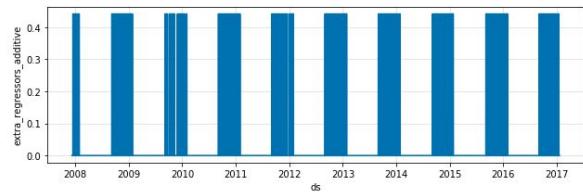
Seasonality

$$s(t)$$



Regressors

$$X(t)$$



Piecewise linear trend

- N changepoints
- Segment-wise independent
- Automatic changepoint detection
- Optional logistic growth

$$g(t) = (k + \mathbf{a}(t)^\top \boldsymbol{\delta})t + (m + \mathbf{a}(t)^\top \boldsymbol{\gamma}),$$

$$a_j(t) = \begin{cases} 1, & \text{if } t \geq s_j \\ 0, & \text{otherwise} \end{cases}$$

$\uparrow \quad \quad \quad \uparrow$
 $a_j(t)$ $-s_j\delta_j$

Seasonality

- N Fourier terms
- Automatic yearly, weekly, daily
- Optional multiplicative mode

$$s(t) = \sum_{n=1}^N \left(a_n \cos \left(\frac{2\pi n t}{P} \right) + b_n \sin \left(\frac{2\pi n t}{P} \right) \right)$$

$$s(t) = X(t)\boldsymbol{\beta}.$$

$$X(t) = \left[\cos \left(\frac{2\pi(1)t}{365.25} \right), \dots, \sin \left(\frac{2\pi(10)t}{365.25} \right) \right]$$

Optional Regularization

$$R(\theta, \epsilon, \alpha) = \frac{1}{p} \sum_{i=1}^p \log \left(\frac{1}{\epsilon \cdot e} + \alpha \cdot |\theta_i| \right) + \log(\epsilon) + 1$$

Events / Holidays

- Automatic for given country
- Various user-defined formats
- Optional multiplicative mode

(Future-Known) Regressors

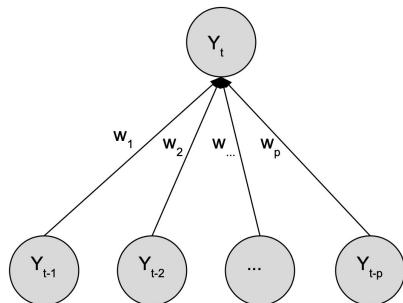
- Single weight
- Real-valued regressor
- Optional multiplicative mode

$$Z(t) = \sum_{i=1}^m c_i e_i(t)$$

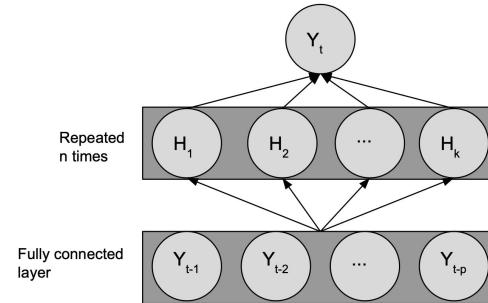
$$R(t) = \sum_{i=1}^l d_i v_i(t)$$

Auto-Regression

- By default AR-Net(0)
- Depth customizable AR-Net(n)
- Optional auto-AR via regularization



AR-Net(0)
Interpretable



AR-Net(n)
Non-linear modeling

(Lagged) Covariates

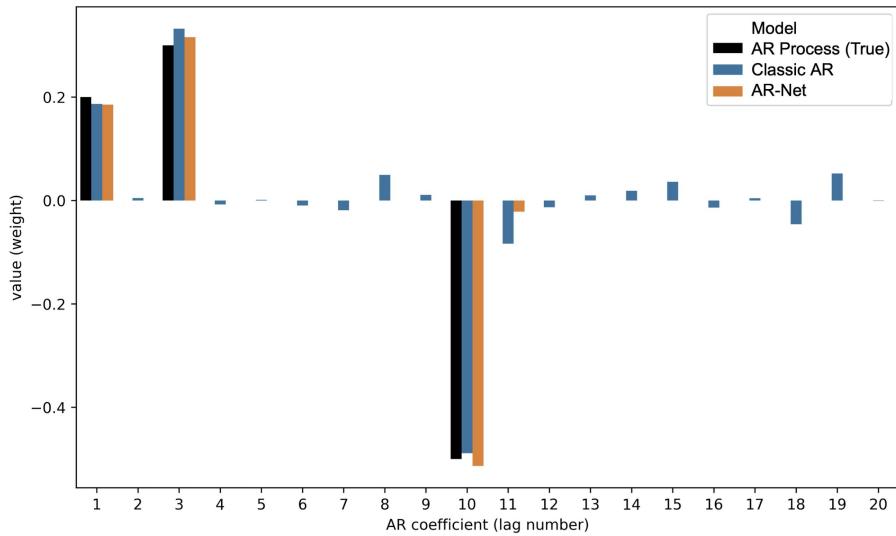
- By default AR-Net(0) with y as target
- Depth customizable AR-Net(n)
- Optional lag-sparsification via regularization

$$\begin{aligned} R_{AR}(\theta) &= R(\theta, \epsilon = 3, \alpha = 1) \\ &= \frac{1}{p} \sum_{i=1}^p \log\left(\frac{1}{3 \cdot e} + |\theta_i|\right) + \log(3) + 1 \end{aligned}$$

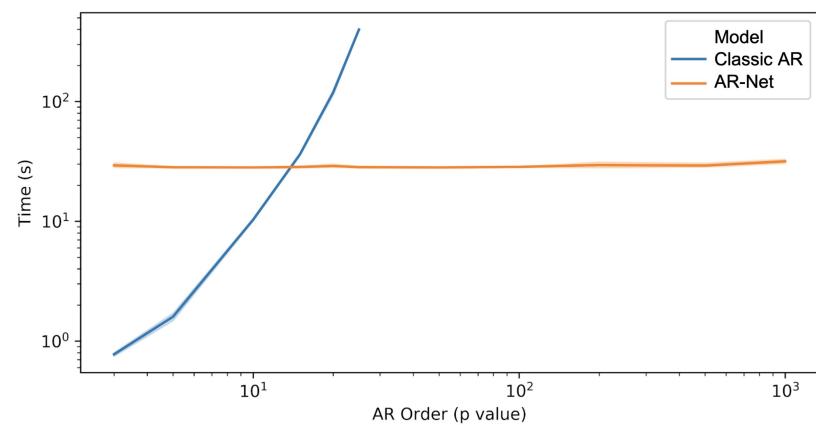
Optional Regularization

Automatic AR-lag selection, yet faster.

Model



Automatic Sparsity



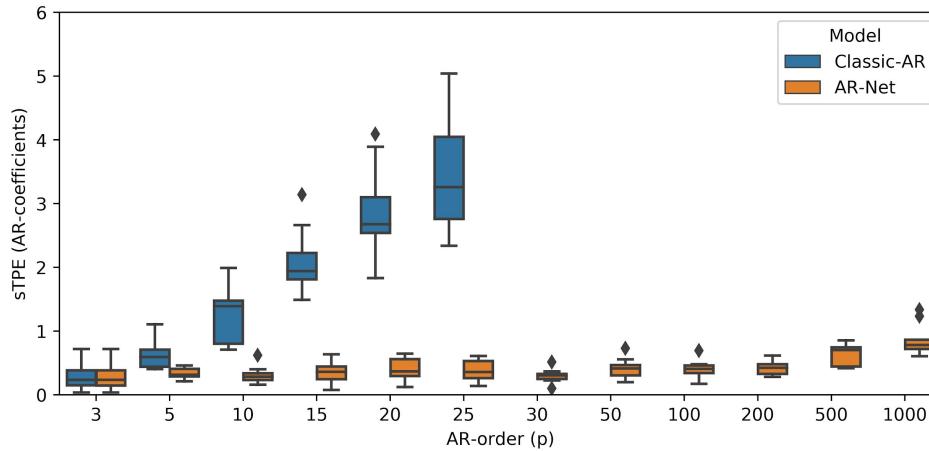
Quadratically faster

Sparse AR-Net surpasses Classic AR and scales to large orders.

Model

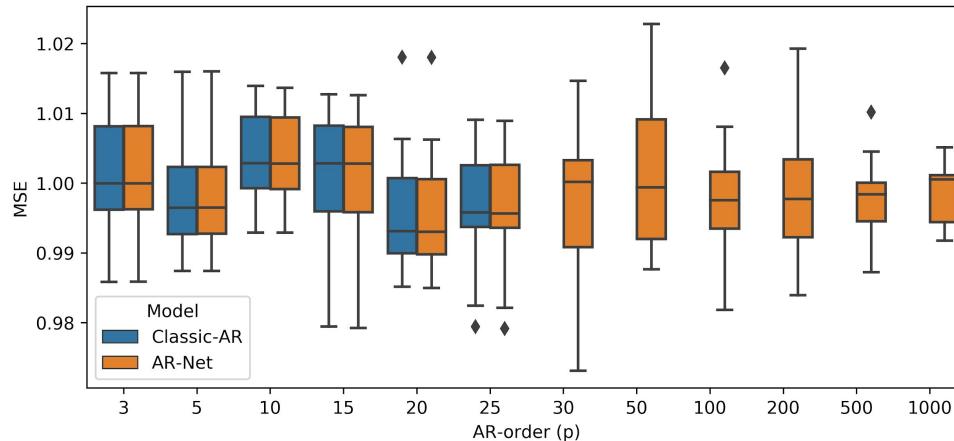
Closeness to true coefficients

$$sTPE = 100 \cdot \frac{\sum_{i=1}^{i=p} |\hat{w}_i - w_i|}{\sum_{i=1}^{i=p} |\hat{w}_i| + |w_i|}$$



MSE loss on forecast target

$$MSE = \frac{1}{n} \sum_1^n (y_t - \hat{y}_t)^2$$



A user-friendly Python package

Gentle learning curve.

Get results first. Learn. Improve.

Powerful, customizable, extendable.

```
m = NeuralProphet()  
metrics = m.fit(df, freq='D')  
forecast = m.predict(df)  
m.plot(forecast)
```

Hyperparameters have smart defaults.

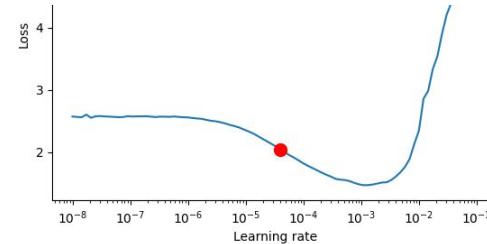
Use

Loss Function is Huber loss,
unless user-defined.

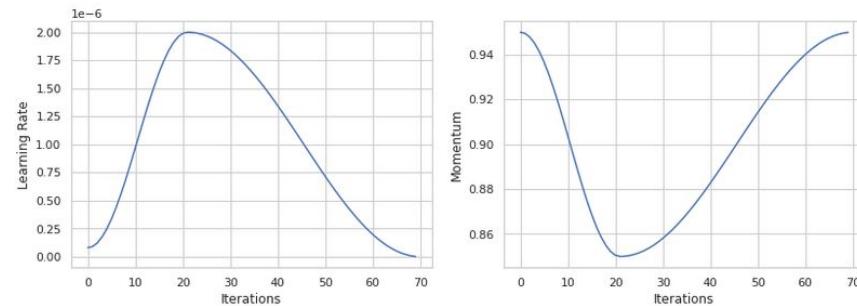
$$L_{\text{huber}}(y, \hat{y}) = \begin{cases} \frac{1}{2\beta}(y - \hat{y})^2, & \text{for } |y - \hat{y}| < \beta \\ |y - \hat{y}| - \frac{\beta}{2}, & \text{otherwise} \end{cases}$$

The learning rate is approximated
with a **learning-rate range test**.

Batch size and epochs are approximated
from the dataset size.



We use **one-cycle policy**
with AdamW as optimizer for simplicity.



Missing Data is automatically filled in:

1. bi-directional linear interpolation
2. centred rolling average

Data is automatically normalized:

Name	Normalization Procedure
'auto'	'minmax' if binary, else 'soft'
'off'	bypasses data normalization
'minmax'	scales the minimum value to 0.0 and the maximum value to 1.0
'standardize'	zero-centers and divides by the standard deviation
'soft'	scales the minimum value to 0.0 and the 95th quantile to 1.0
'soft1'	scales the minimum value to 0.1 and the 90th quantile to 0.9

We have utils ...

Visualize:

- Plot past and future predictions
- Decompose forecast components
- Interpret model parameters
- Plot most recent prediction
- Inspect a particular forecast horizon

Other:

- Simple Split, cross validation, double cross validation
- Control logger verbosity
- Make fit reproducible
- Global modelling
- Benchmarking
- ...

Quick Start Examples

The classic: Temperature forecast

Using Trend, Seasonality and Auto-Regression:

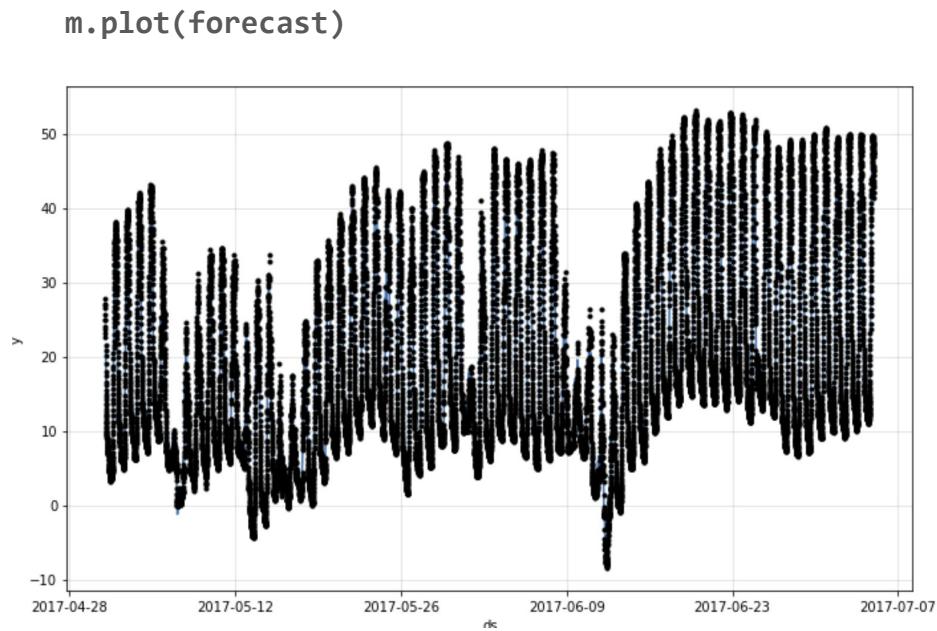
Example 1: 1 step ahead

Example 2: 36 steps ahead

Dataset:

Observed temperature in Yosemite Valley, measured every 5 min over two months.

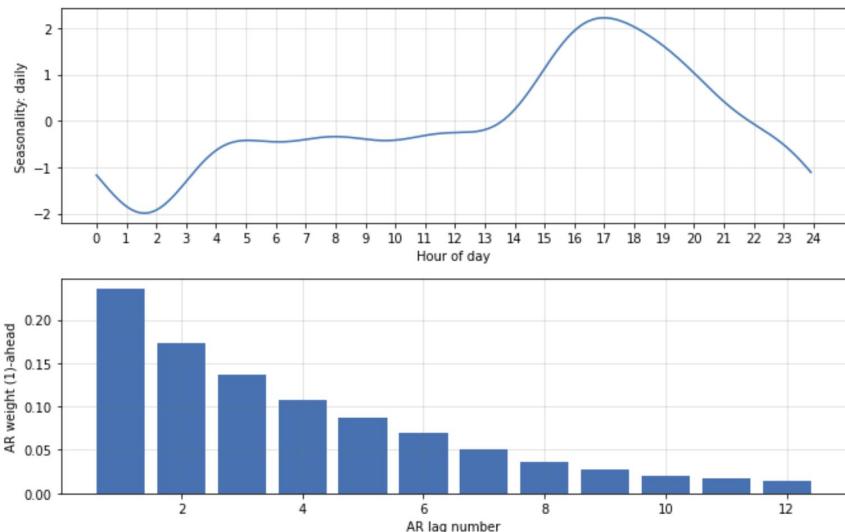
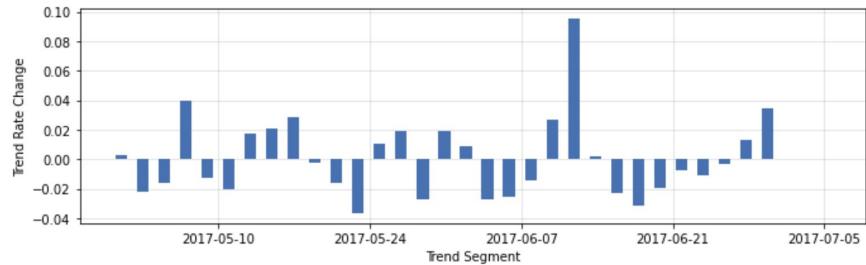
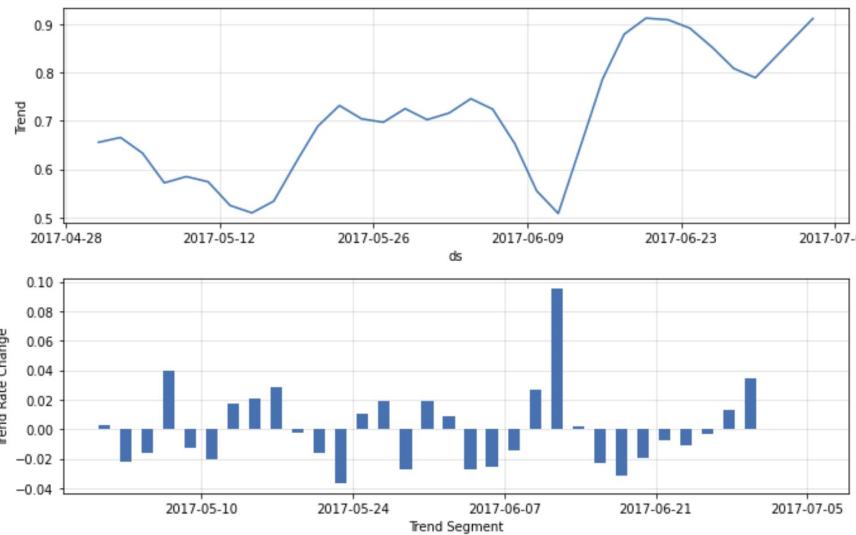
```
m = NeuralProphet(n_lags=12)
metrics = m.fit(df, freq='D')
forecast = m.predict(df)
```



Visualize model parameters with one line.

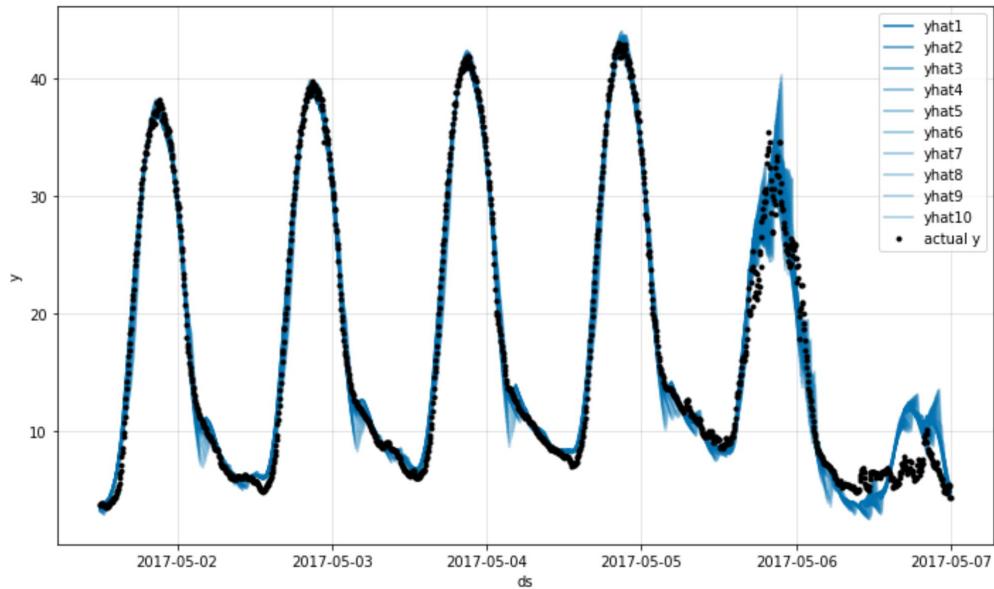
Example: Yos1

`m.plot_parameters()`



```
m = NeuralProphet(  
    n_lags=72,  
    n_forecasts=36,  
)  
  
metrics = m.fit(df, freq='D')  
  
forecast = m.predict(df)
```

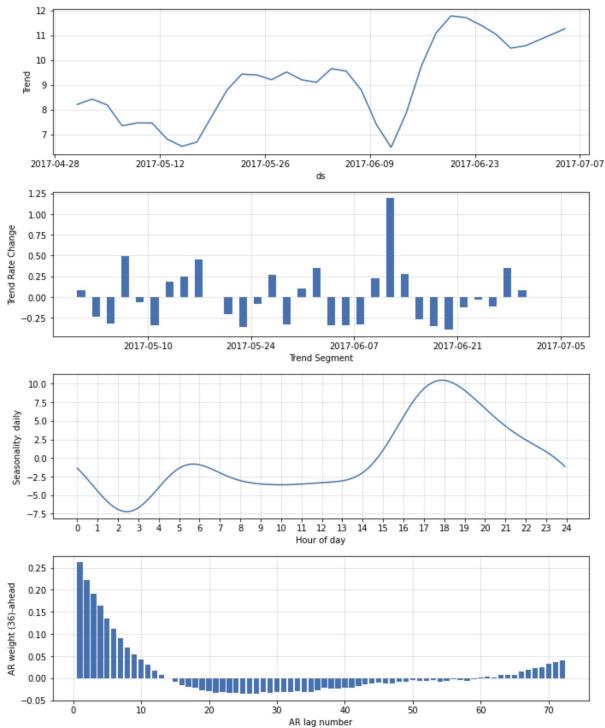
```
m.plot(forecast[:6*24*12])
```



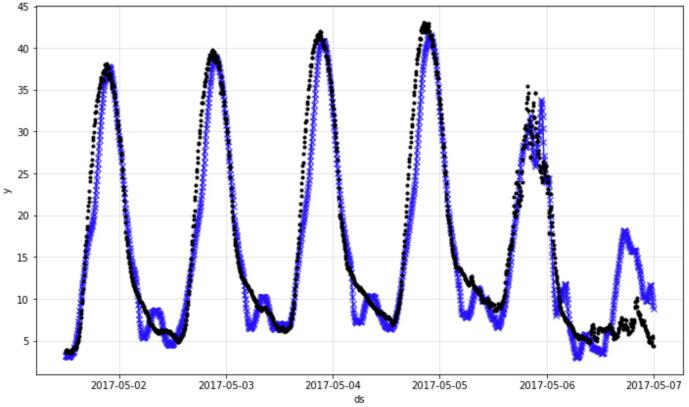
Analyze a specific forecast horizon. Sparsify.

Example: Yos36

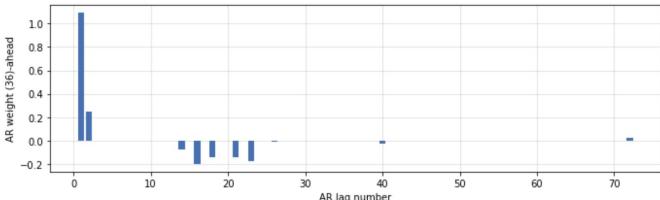
```
m.highlight_nth_step_ahead_of_each_forecast(36)  
m.plot_parameters()
```



```
m.plot(forecast[:6*24*12])
```



NeuralProphet(ar_sparsity=0.1, ...)

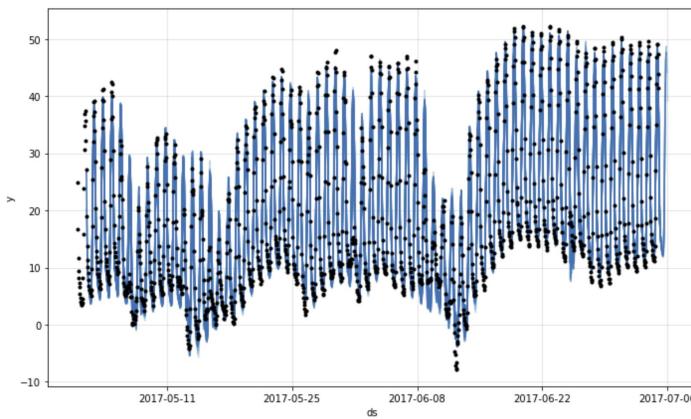


Want to forecast a larger horizon?

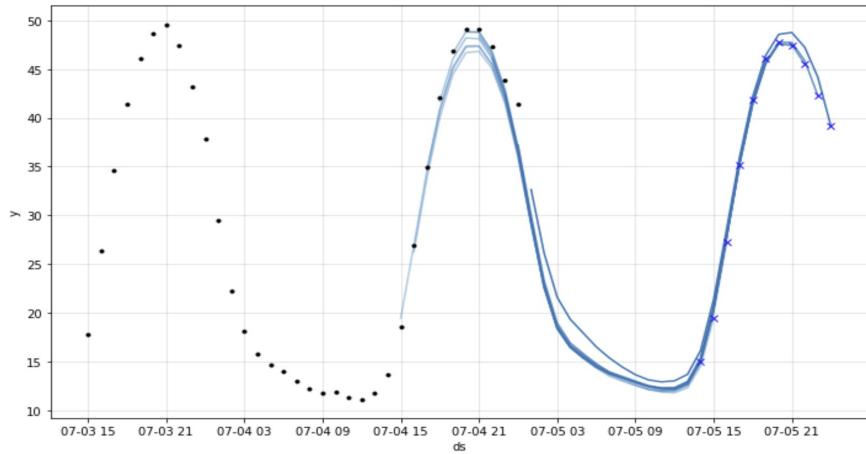
Example: Yos24

```
df_hourly = df.set_index('ds', drop=False).resample('H').mean().reset_index()
```

```
m = NeuralProphet(  
    n_lags=24,  
    n_forecasts=24,  
)  
metrics = m.fit(df_hourly, freq='H')  
forecast = m.predict(df_hourly)  
m.plot(forecast)
```



```
future = m.make_future_dataframe(df_hourly)  
forecast = m.predict(future)  
m = m.highlight_nth_step_ahead_of_each_forecast(24)  
m.plot_last_forecast(forecast, include_previous_forecasts=10)
```



We are working hard to extend the framework. Join us!

Extensions [upcoming]

- Hierarchical Forecasting & Global Modelling
- Quantifiable and Explainable Uncertainty
- Anomaly Prediction & Semi-Supervised Learning
- Attention: Automatic Multimodality & Dynamic Feature Importance

Improvements [upcoming]

- Improved NN
- Faster Training Time & GPU support
- Improved UI
- Diagnostic Tools for Deep Dives

Extension: Global Modelling

Scenario: Dataset with many timeseries of the same domain.

Goal: Fit a single forecasting model with shared weights

Why:

- Better generalization and model size savings
- A single time series may not reflect the entire time series dynamics
 - e.g. anomaly prediction (Zhu and Laptev 2017), tourism demand (Claveria and Torra 2014)
- Extrapolation to new times series

Example Demonstrations:

- Energy load: Forecasting on previously unseen data
- Engine faults: Predicting anomalies with a single model

Forecast energy load of region without prior data.

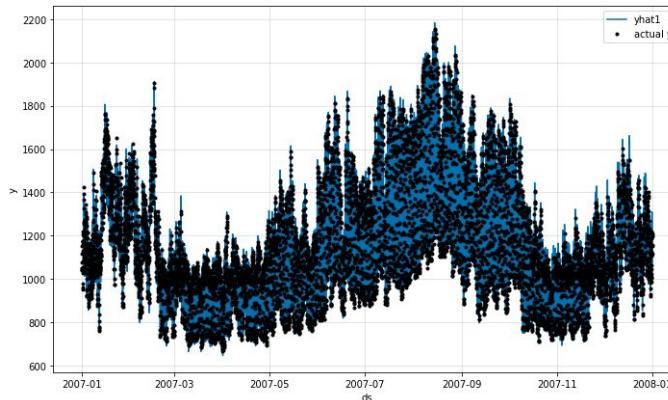
ERCOT hourly load

Data: Train on 3 years, 1 year test (unseen region North)

Model: 24 hours lags, 1 step ahead

```
Coast  
East  
...  
West  
(Without  
North)  
  
North → m.test(df_north)
```

→ `m = NeuralProphet(n_lags=24)`
 `m.fit(list_of_dfs,freq='H')`



Model	Train dataset	MAE	RMSE
1. Global NP	All regions but North	27.45	33.95
2. Local NP	North	16.09	21.47
3. Local Prophet	North	147.30	191.58
4. Naive Model	North	175.58	230.09

Engine failure - supervised class prediction

Data: NASA ProDiMES Simulator

- 500 engines with 40 flights each (time-stamps)
- Test on unseen data: 100 engines

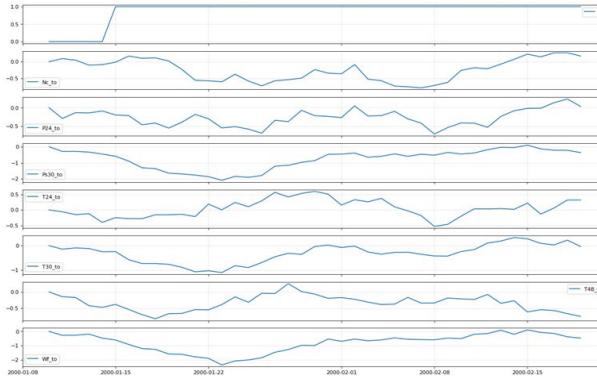
Model:

- 1 step ahead (predict class of next flight)
- Covariates: 5 lags from 14 sensors

```

-----| → m = Classification_NP()
      | Engine 1   m = m.add_lagged_regressor(
      |           |           names=sensors,
      |           |           n_regressors=5)
      |           | m.fit(list_of_dfs, freq='D')
      |
      | ...
      | Engine 400
      |
      | ...
      | Engine 401 → m.test(list_of_test_dfs)
      |
      | ...
      | Engine 402
      |
      | ...
      | Engine 500
  
```

sensors



Model	Accuracy	Balanced Acc
1. Global Classifier NP	0.79	0.80
2. Logistic Regression	0.77	0.75
3. Dummy Classifier	0.57	0.50

Global Modelling

- Local normalization
- Local/global model components
 - E.g. Trend, seasonality

Class Prediction

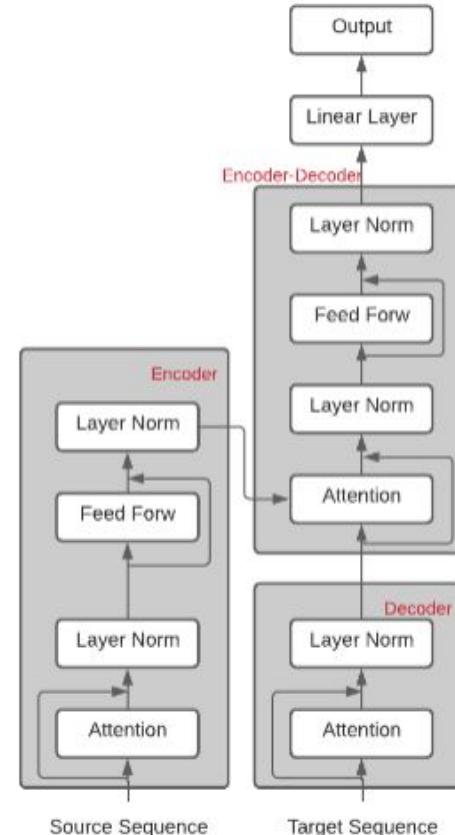
- Cross-correlate covariates
- Multi-class targets
- Semi-supervised approach w. Attention

Extension: Attention Dynamic Weighted Features & Ensembles

- Transformer models have become popular in the world of NLP over the last few years.
- Transformers get its power from its attention module.
 - Attention has the ability to capture relationships between each feature with every other feature within the sample.

Transformer Models

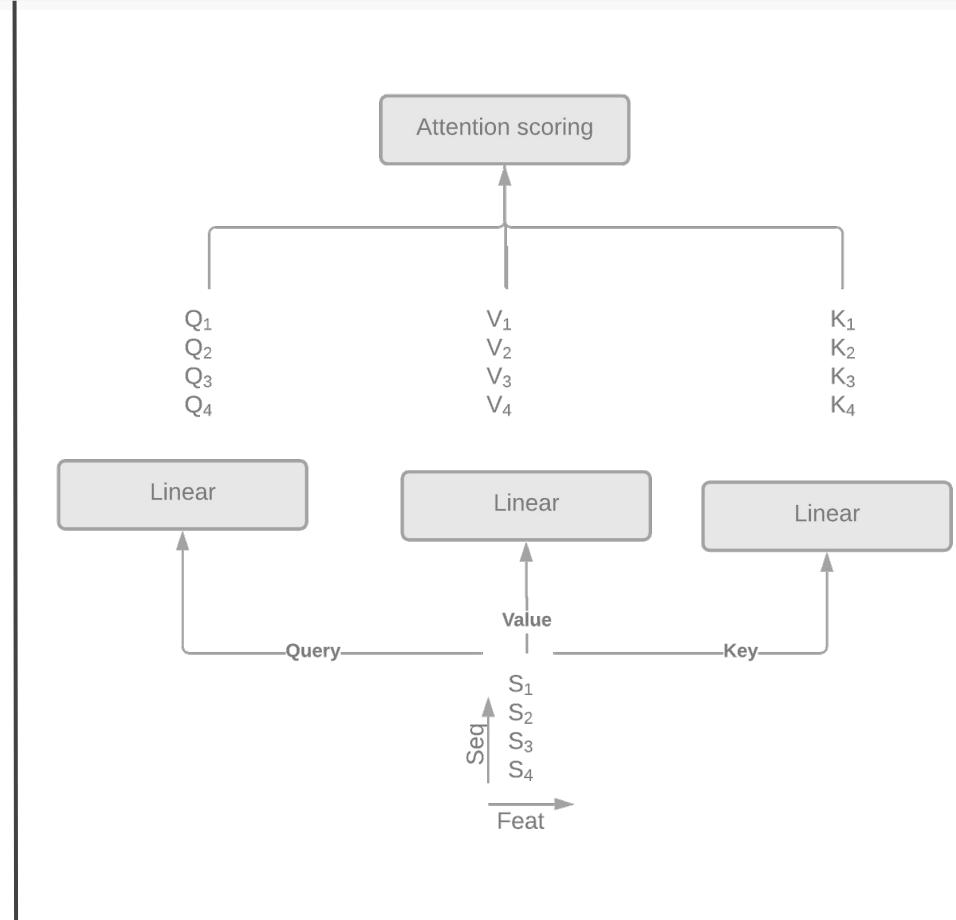
- **Encoder:** Pays self-attention over source sequence.
- **Decoder:** Pays self-attention over Target sequence.
- **Encoder-Decoder:** Target pays attention to source.



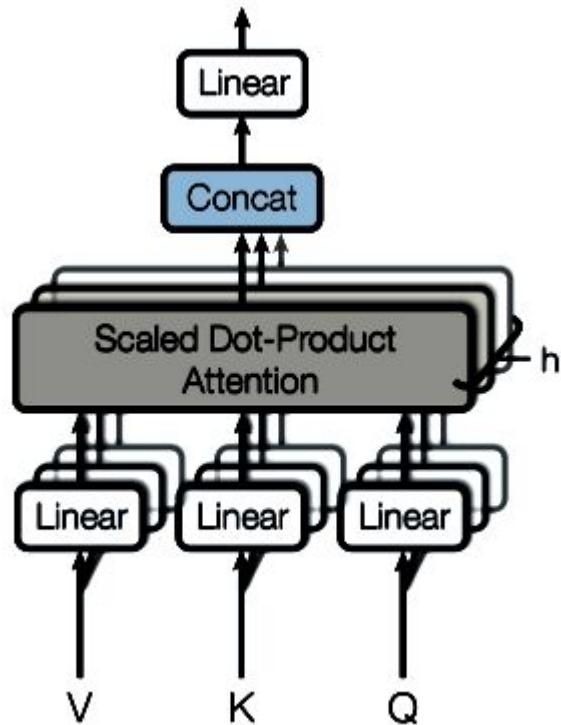
Single Attention

Dot Product Attention
Scoring

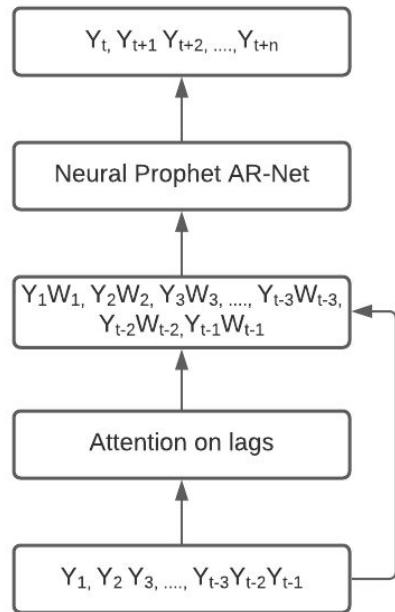
$$\begin{aligned} & Q_1 K_1 V_1 + Q_1 K_2 V_2 + Q_1 K_3 V_3 + Q_1 K_4 V_4 \\ & Q_2 K_1 V_1 + Q_2 K_2 V_2 + Q_2 K_3 V_3 + Q_2 K_4 V_4 \\ & Q_3 K_1 V_1 + Q_3 K_2 V_2 + Q_3 K_3 V_3 + Q_3 K_4 V_4 \\ & Q_4 K_1 V_1 + Q_4 K_2 V_2 + Q_4 K_3 V_3 + Q_4 K_4 V_4 \end{aligned}$$



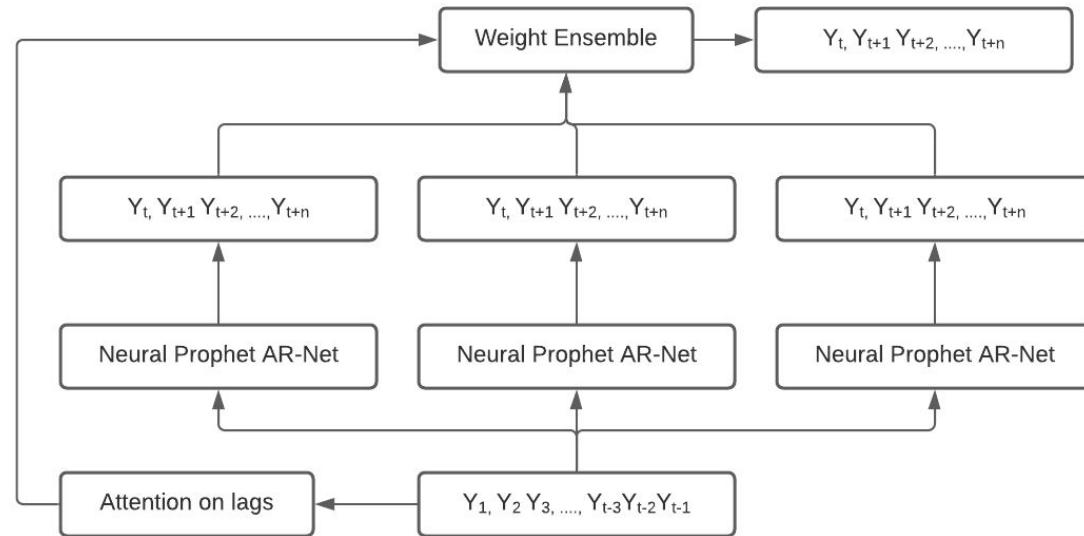
Multihead Attention



Dynamic Feature Weighting



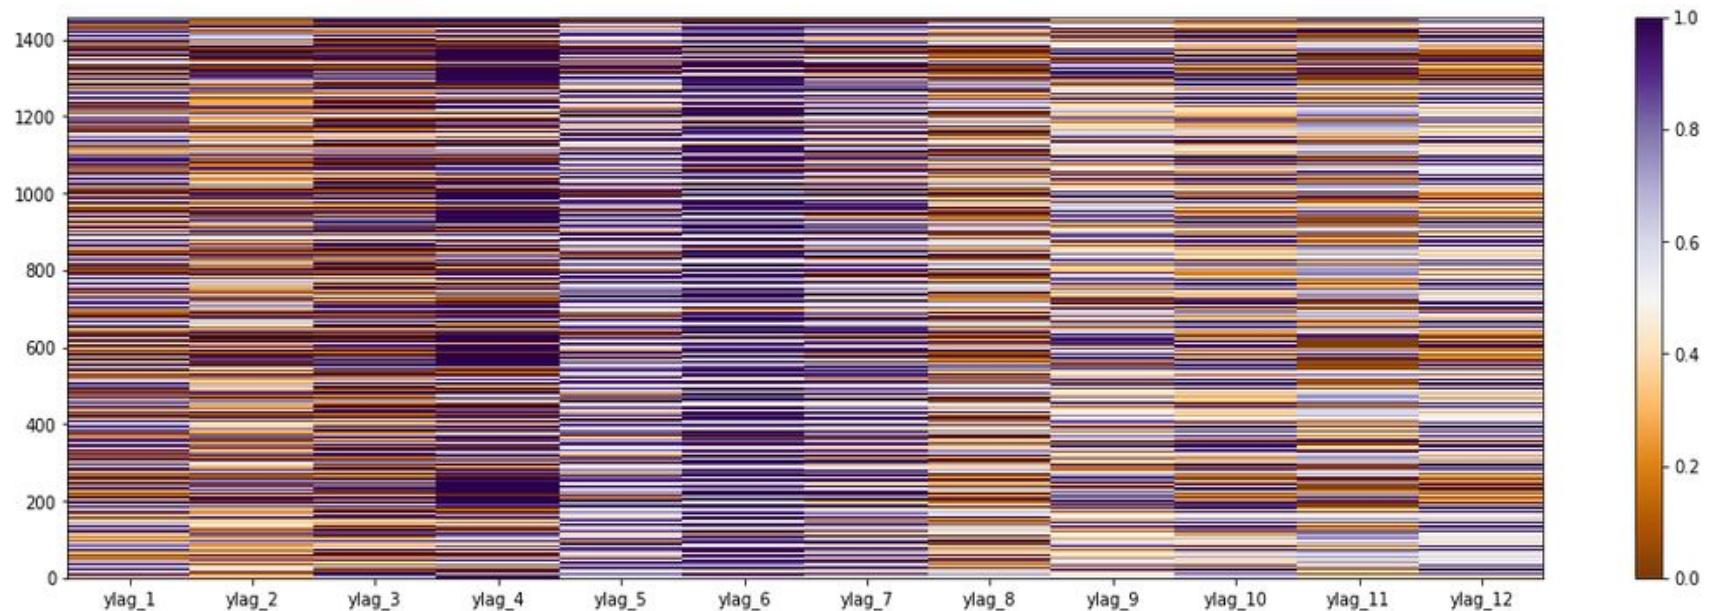
Dynamic Ensemble Weighting



Preliminary Results (MASE)

Dataset	Model	Horizon		
		6	12	24
Solar-Irradiance	w/ Attention	1.087	17.686	18.756
	w/o Attention	2.69	19.46	24.714
ETT-h1	w/ Attention	35.995	92.373	186.349
	w/o Attention	36.761	92.695	188.281
Energy Load	w/ Attention	62.246	144.266	259.45
	w/o Attention	65.423	152.005	274.744

Feature Importance Plot



Ablation studies planned

- Splitting of multiple heads in Multi-head attention
 - One head per lag/feature variable
 - Strategic Feature split
 - Copy instead of split
- Combining of multiple heads in Multi-head attention
 - Sum→sigmoid
 - Concatenate→Linear Layer→sigmoid
 - sum→sigmoid/len(attention out)
 - Concatenate→Linear Layer→sigmoid/len(attention out)

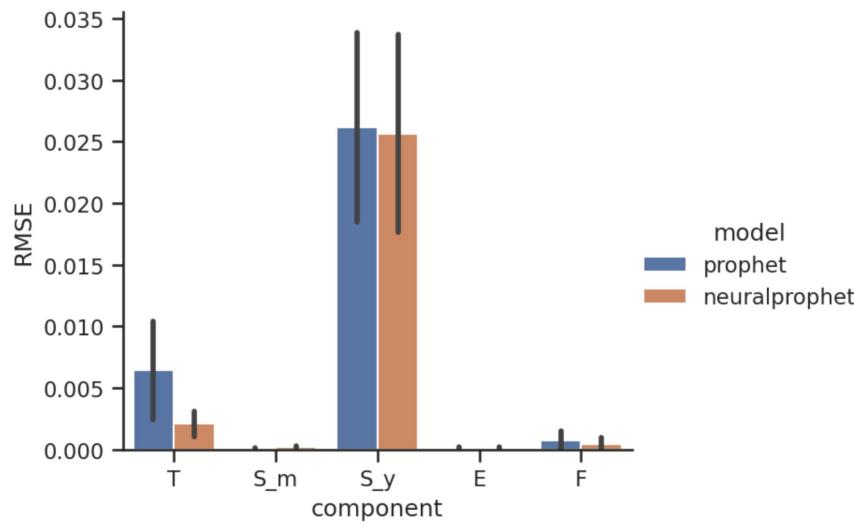
- Implementing attention on neural prophet's AR-Net model has not only made the model more explainable but also improved the forecast accuracy on univariate models.
- Future directions will be to extend this for global and multivariate models.

NeuralProphet is a modern Prophet with a superset of its features.

Task	Prophet	NeuralProphet
Very small dataset (less than 100 samples)	✓	
Large dataset (more than 1000 samples)		✓
Long range forecast (e.g. multiple years)	✓	✓
Short to medium range forecast (e.g. 1 to 1000 steps ahead)		✓
Specific forecast horizon (e.g. next 24h)		✓
Auto-correlation (dependence on previous observations)		✓
Lagged regressors (observed covariates)		✓
Non-linear dynamics		✓
Global modelling of panel dataset		✓
Fast prediction (computational inference time)		✓

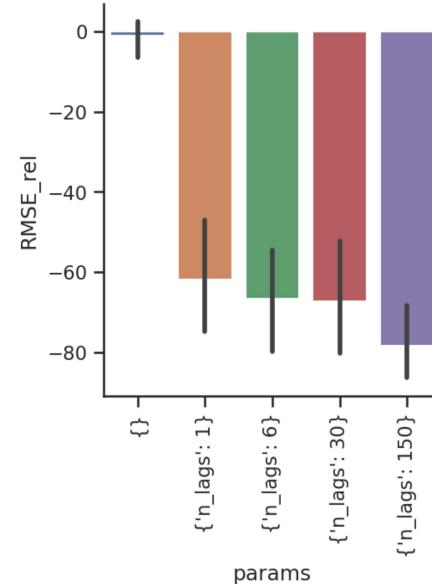
Accuracy of forecast decomposition

On synthetic dataset.



Forecast error difference to Prophet

On a collection of real datasets.



Simple and powerful forecaster,
without compromising on interpretability.

1:1 replacement for Prophet,
with many new capabilities, superior for most applications.

Nothing but Python & PyTorch,
extensible to future state-of-the-art in forecasting.

THANK YOU, dear collaborator, supporter and advisor!

Team



facebook



Skoltech



Oskar Triebe



Hansika Hewamalage



Mateus De Castro Ribeiro



Lluvia Ochoa



Nikolay Laptev



Polina Pilyugina

Abishek Sriramulu

Ram Rajagopal

Christoph Bergmeir

Alessandro Panella

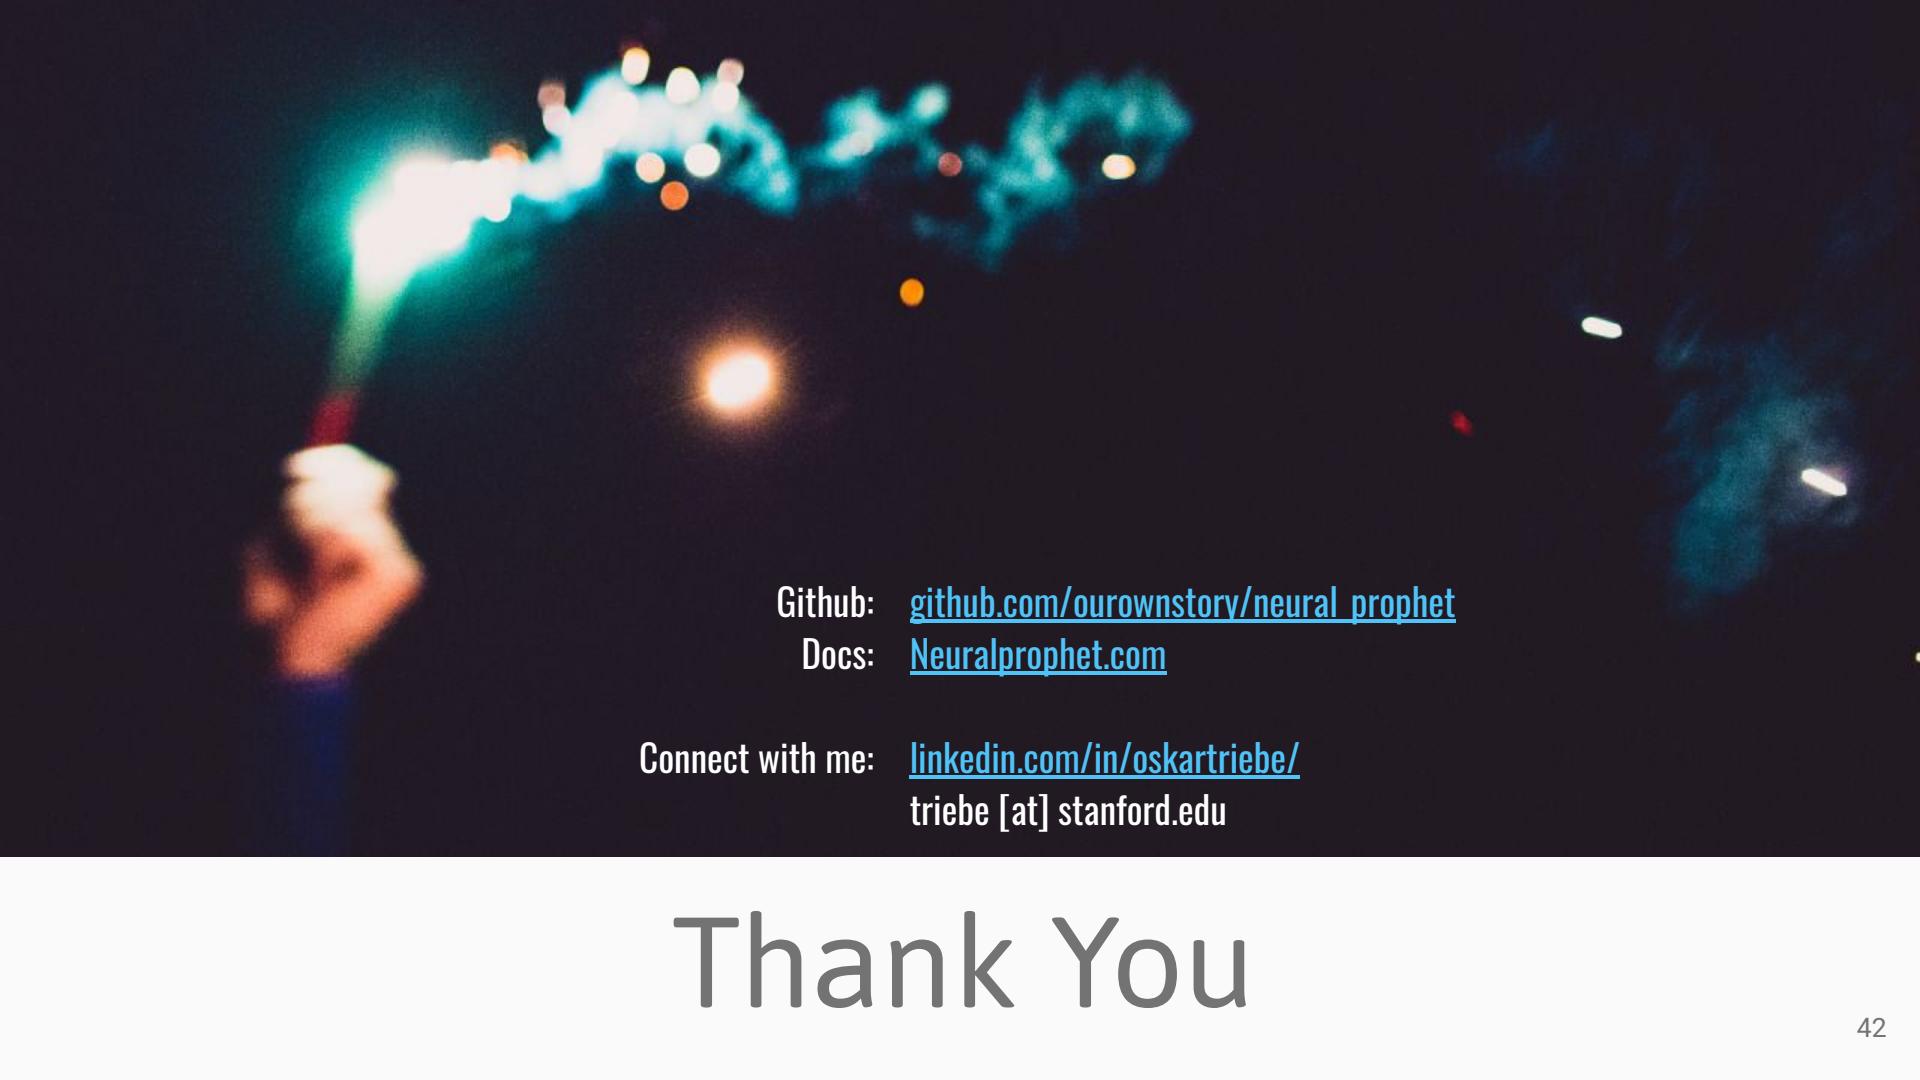
Evgeny Burnaev

Caner Komurlu

Italo Lima

Gonzague Henri

Bernhard Hausleitner



Github: github.com/ourownstory/neural_prophet
Docs: Neuralprophet.com

Connect with me: linkedin.com/in/oskartriebe/
triebe [at] stanford.edu

Thank You

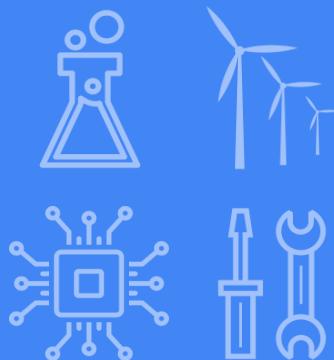
References

- Claveria, O., & Torra, S. (2014). Forecasting tourism demand to Catalonia: Neural networks vs. time series models. *Economic Modelling*, 36, 220-228.
- Zhu, L., & Laptev, N. (2017, November). Deep and confident prediction for time series at uber. In *2017 IEEE International Conference on Data Mining Workshops (ICDMW)* (pp. 103-110). IEEE.

Appendix: AR-Net

A time series is

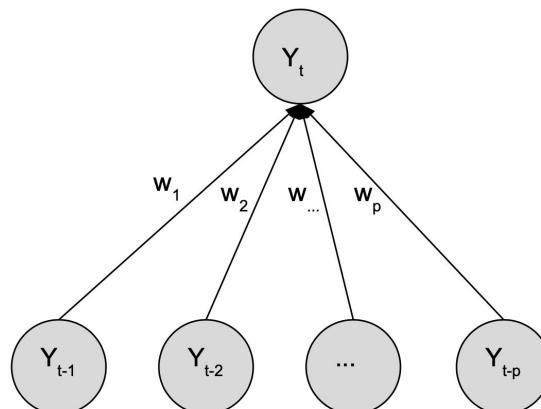
a sequence of data points
that occur in successive order
over some period of time.



AR-Net is a Neural Network for autoregressive time series.

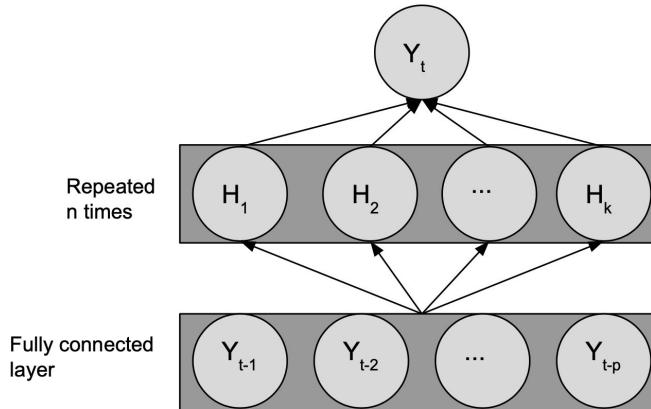
AR-Net(0)

Interpretable



AR-Net(n)

Stronger modeling ability



Skip estimating the AR process order.

Model: AR-Net

$$y_t = c + \sum_{i=1}^{i=p} w_i * y_{t-i} + e_t$$



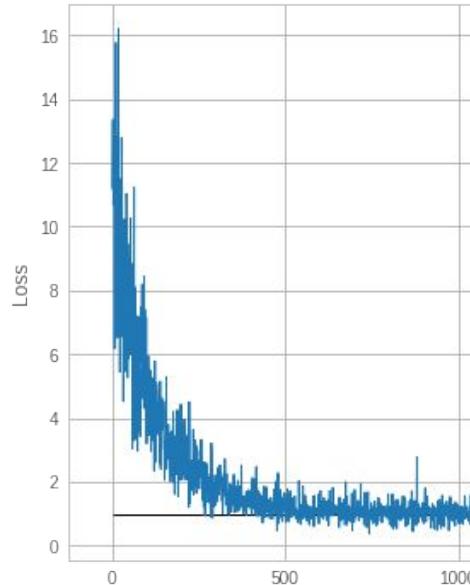
$$\min_{\theta} L(y, \hat{y}, \theta) + \lambda(s) \cdot R(\theta)$$

$$\lambda(s) = c_\lambda \cdot (s^{-1} - 1)$$

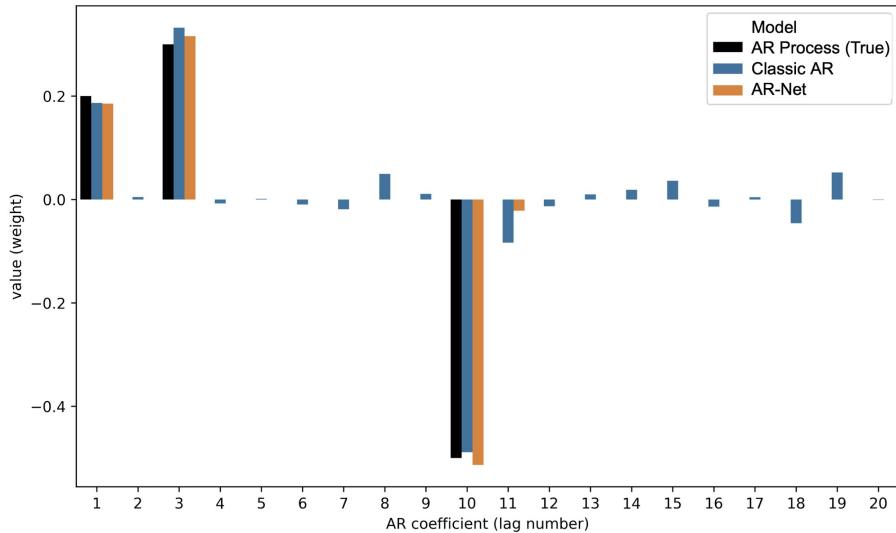
$$s = \frac{\hat{p}_{data}}{p_{model}}$$

$$c_\lambda \approx \frac{\sqrt{\hat{L}}}{100}$$

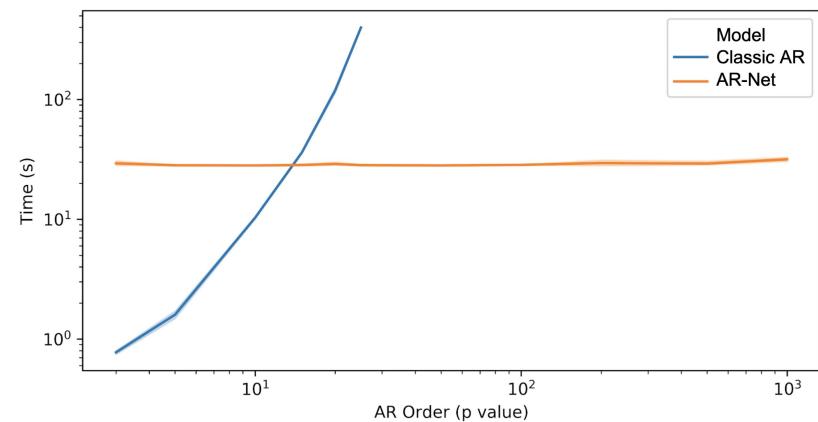
$$R(\theta) = \frac{1}{p} \sum_{i=1}^p \frac{2}{1 + \exp(-c_1 \cdot |\theta_i|^{\frac{1}{c_2}})} - 1$$



Trained with SGD (Adam)



Automatic Sparsity



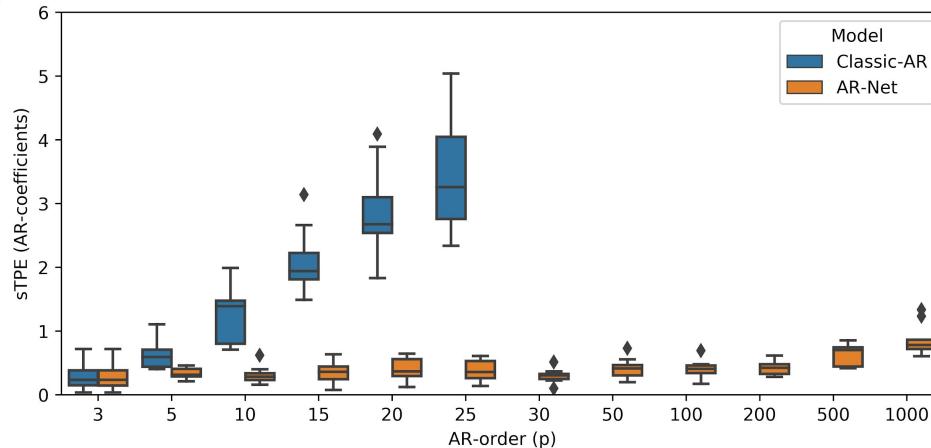
Quadratically faster

Sparse AR-Net surpasses Classic AR and scales to large orders.

Model: AR-Net

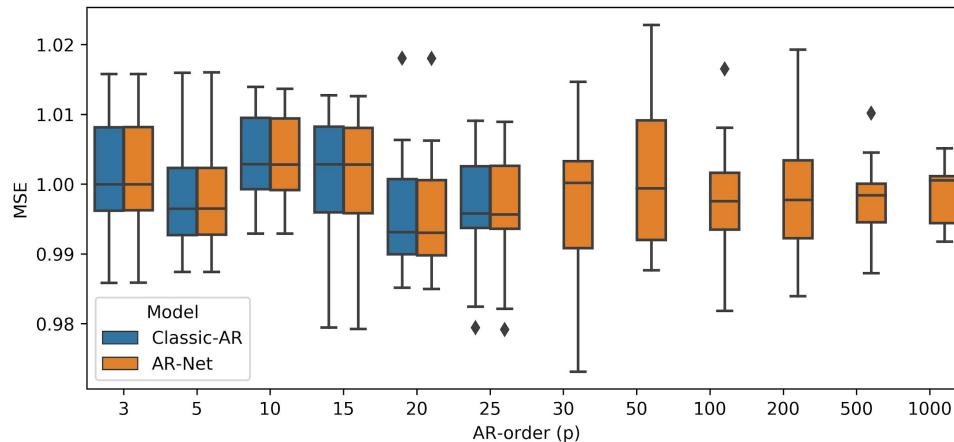
Closeness to true coefficients

$$sTPE = 100 \cdot \frac{\sum_{i=1}^{i=p} |\hat{w}_i - w_i|}{\sum_{i=1}^{i=p} |\hat{w}_i| + |w_i|}$$



MSE loss on forecast target

$$MSE = \frac{1}{n} \sum_1^n (y_t - \hat{y}_t)^2$$



Appendix: Model Use Details

skip data preparation:

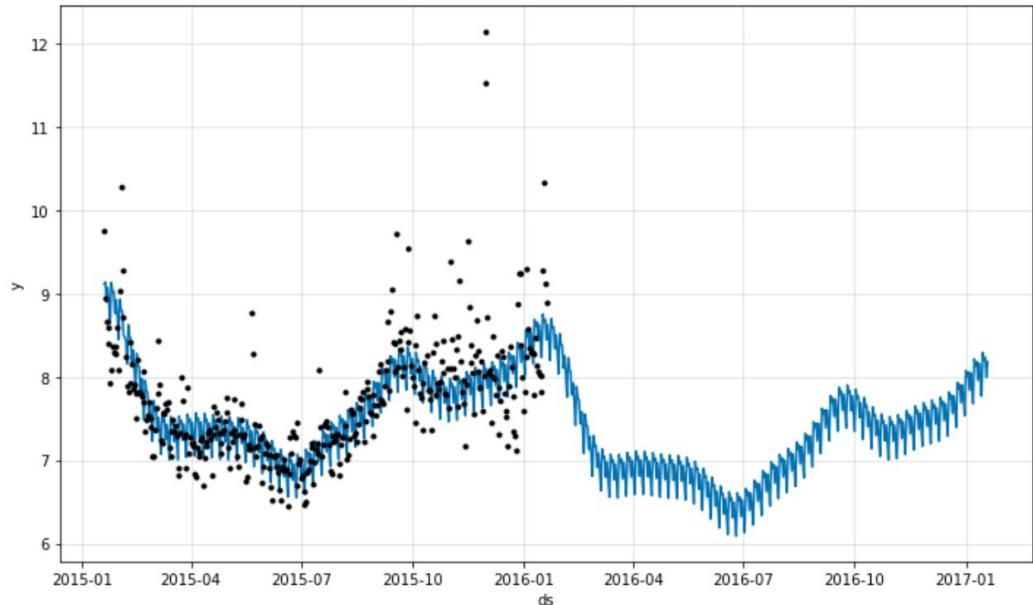


Just create a DataFrame
with desired columns

```
import pandas as pd
from neuralprophet.neural_prophet import NeuralProphet

df = pd.read_csv('..../data/example_wp_log_peyton_manning.csv')
```

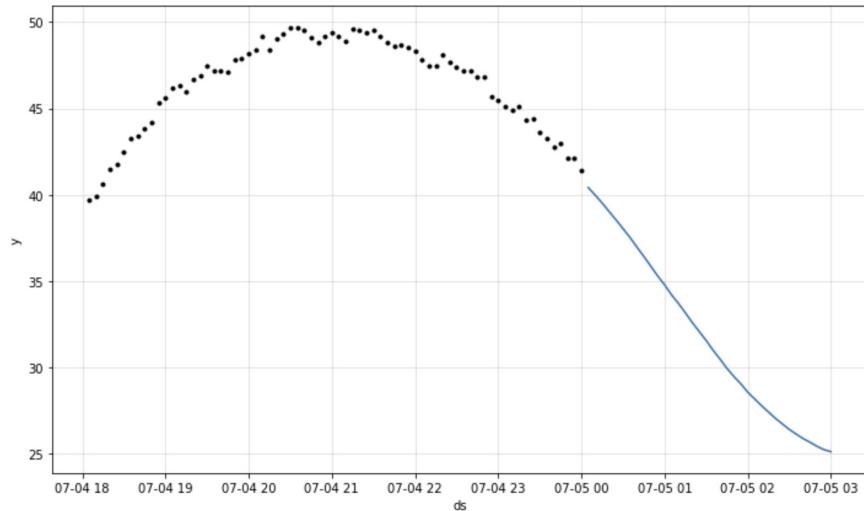
```
# linear time-dependent model
m = NeuralProphet()
metrics = m.fit(df)
future = m.make_future_dataframe(df, future_periods=365)
forecast = m.predict(future)
fig_fcst = m.plot(forecast[-730:])
```



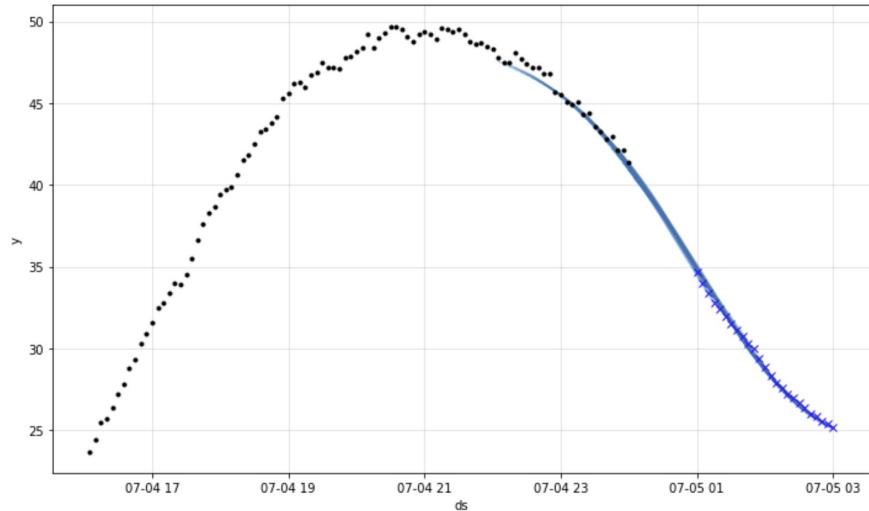
Predict. See how new data impacts the forecast.

Example: Yos36

```
m = m.highlight_nth_step_ahead_of_each_forecast(None) # reset highlight  
fig = m.plot_last_forecast(forecast)
```



```
m = m.highlight_nth_step_ahead_of_each_forecast(3*12)  
fig = m.plot_last_forecast(forecast, include_previous_forecasts=2*12)
```



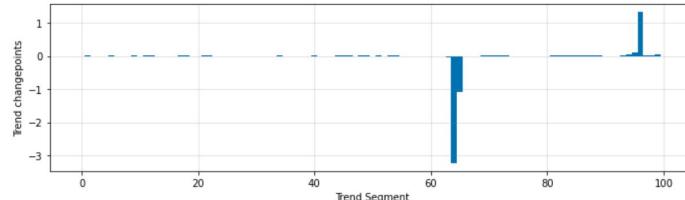
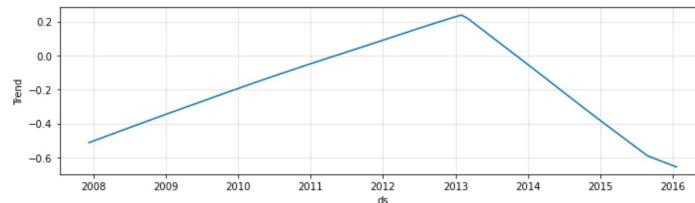
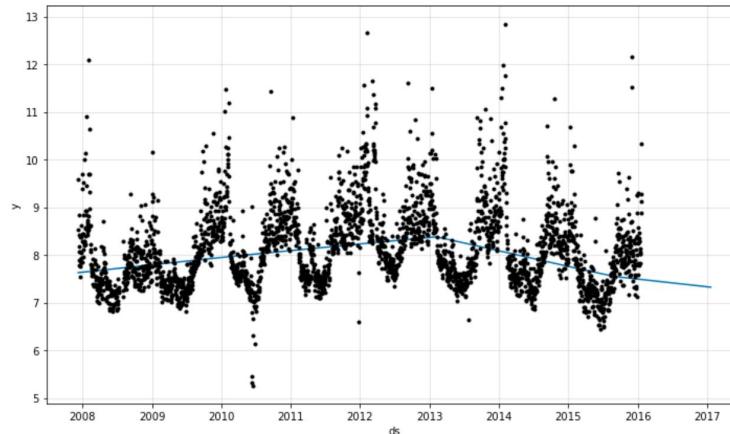
```
# or evaluate while training
m = NeuralProphet()
metrics = m.fit(df, validate_each_epoch=True, valid_p=0.2)
metrics.tail()
```

Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.

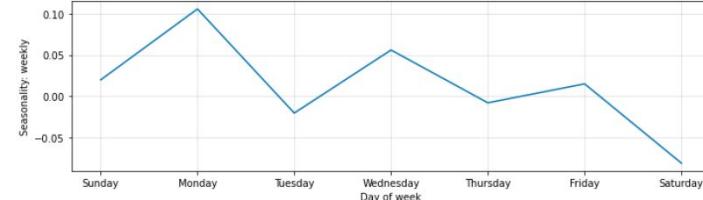
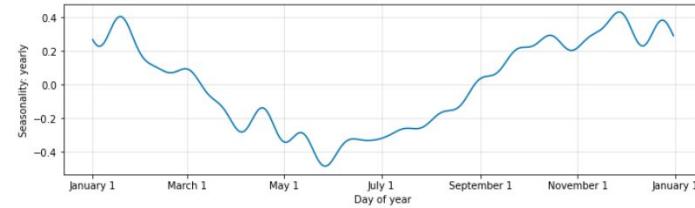
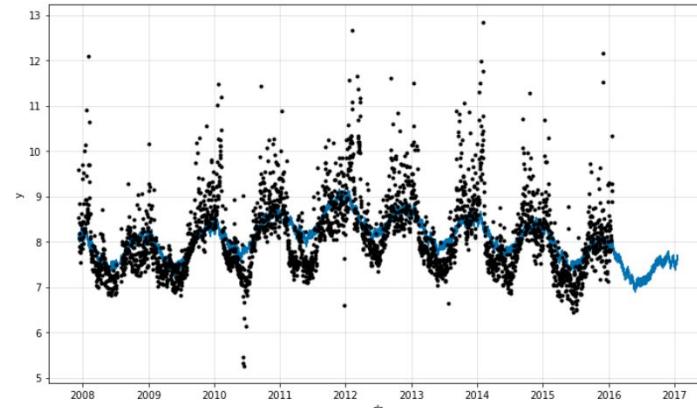
	SmoothL1Loss	MAE	RegLoss	SmoothL1Loss_val	MAE_val
35	0.163102	0.371323	0.0	0.485371	0.779465
36	0.161851	0.369609	0.0	0.368921	0.648736
37	0.161122	0.369219	0.0	0.366328	0.648230
38	0.168598	0.376638	0.0	0.348269	0.627886
39	0.167961	0.375777	0.0	0.362161	0.642699

```
# split manually
m = NeuralProphet()
df_train, df_val = m.split_df(df, valid_p=0.2)
train_metrics = m.fit(df_train)
val_metrics = m.test(df_val)
```

```
m = NeuralProphet(  
    n_changepoints=100,  
    trend_smoothness=2,  
    yearly_seasonality=False,  
    weekly_seasonality=False,  
    daily_seasonality=False,  
)  
metrics = m.fit(df)
```



```
m = NeuralProphet(  
    yearly_seasonality=16,  
    weekly_seasonality=8,  
    daily_seasonality=False,  
    seasonality_reg=1,  
)  
metrics = m.fit(df)
```



Events can be added in different forms.

Hands-on

```
superbowls = pd.DataFrame({'event': 'superbowl',
 'ds': pd.to_datetime(['2010-02-07', '2014-02-02', '2016-02-07'])})
```

```
events_df = pd.concat((superbowls, playoffs))
```

```
m = NeuralProphet()
```

```
m = m.add_country_holidays("US")
```

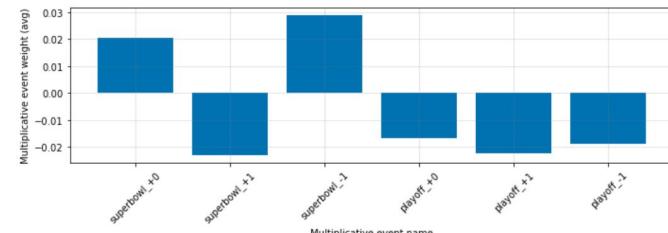
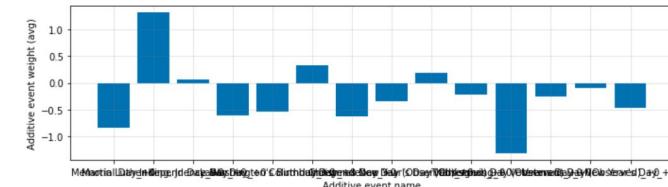
```
m = m.add_events(
    ["superbowl", "playoff"],
    lower_window=-1,
    upper_window=1,
    mode="multiplicative",
    regularization=0.5
)
```

```
history_df = m.create_df_with_events(df, events_df)
metrics = m.fit(history_df)
```

Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.

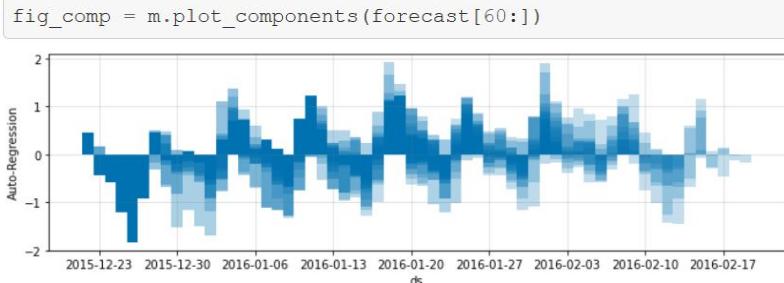
```
future = m.make_future_dataframe(
    history_df,
    events_df,
    future_periods=30,
    n_historic_predictions=30
)
```

```
forecast = m.predict(future)
```



```
m = NeuralProphet(  
    n_forecasts=30,  
    n_lags=60,  
    ar_sparsity=0.1,  
    yearly_seasonality=False,  
    weekly_seasonality=False,  
    daily_seasonality=False,  
)  
  
metrics = m.fit(df)  
future = m.make_future_dataframe(df, n_historic_predictions=30)  
forecast = m.predict(future)
```

Autoregression,
here with sparsity



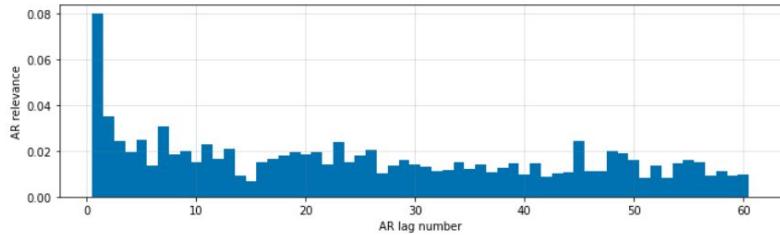
Recommended example notebook:

https://github.com/ourownstory/neural_prophet/blob/master/example_notebooks/autoregression_yosemite_temps.ipynb

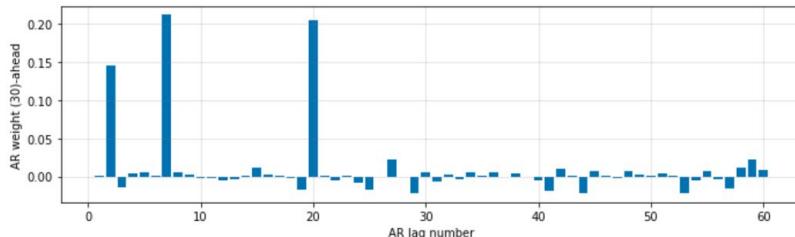
Focus on a specific forecast.

Hands-on

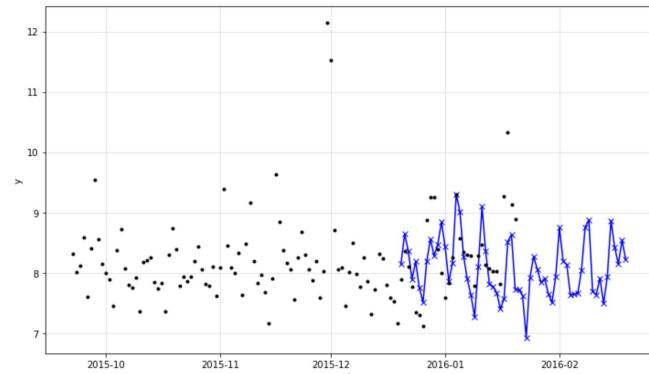
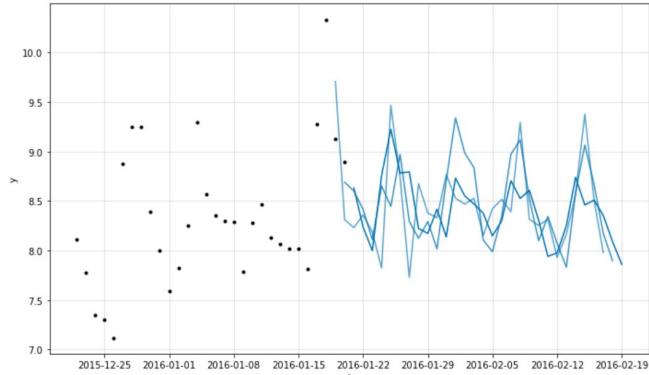
```
fig_param = m.plot_parameters()
```



```
fig_param30 = m.highlight_nth_step_ahead_of_each_forecast(30).plot_parameters()
```



```
fig_prediction = m.plot_last_forecast(forecast[60:],  
include_previous_forecasts=2)
```



Lagged Covariates & Future Regressors

Hands-on

```
m = NeuralProphet(  
    n_forecasts=30,  
    n_lags=60,  
    ar_sparsity=0.1,  
    yearly_seasonality=False,  
    weekly_seasonality=False,  
    daily_seasonality=False,  
)
```

```
df = pd.read_csv('../data/example_wp_log_peyton_manning.csv')  
df['A'] = df['y'].rolling(7, min_periods=1).mean()  
df['B'] = df['y'].rolling(60, min_periods=1).mean()  
m = m.add_covariate(name='A')  
m = m.add_regressor(name='B')
```

```
metrics = m.fit(df)  
future = m.make_future_dataframe(df, n_historic_predictions=30)  
forecast = m.predict(future)
```

Covariates and regressors are similarly added

Make it non-linear.

```
m = NeuralProphet(  
    n_forecasts=30,  
    n_lags=60,  
    learning_rate=1.0,  
    loss_func='Huber',  
    normalize_y=True,  
    num_hidden_layers=2,  
    d_hidden=64,  
)
```

```
fig_comp = m.plot_components(forecast[60:])
```

