

Neural Prophet

Extendable and Scalable Forecasting

Oskar Triebe

Stanford University, October 26, 2020

40th International Symposium on Forecasting

Motivation

Introduction

Model

Hands-On

Outlook

Forecasting Landscape

Traditional Methods

(S)ARIMA(X)

(V)ARMA(X)

GARCH

(S)Naïve

Gaussian Process

HMM

(T)BATS

Seasonal + Trend Decomposition

Exponential Smoothing

Holt-Winters

Dynamic Linear Models

Prophet

AR-Net

ES-RNN

LSTM

N-BEATS

Transformer

WaveNet

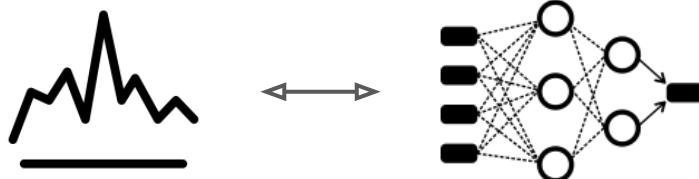
CNN

Other ML

Deep Learning

Chasm

Time series - applied ML



Need expertise in both
domains

Bridge

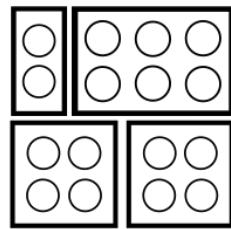
NeuralProphet



Abstracts time series and
ML knowledge

Introduction: NeuralProphet

- Beginner-friendly
- Customizable model
- Interpretable
- Scalable to large data
- Extendable code
- Workflow integration



Add more data, features,
or modelling blocks



Explain model parameters,
Visualize prediction components

The original Prophet introduced Analyst-in-the-loop.

Introduction: Prophet



Forecasting at Scale

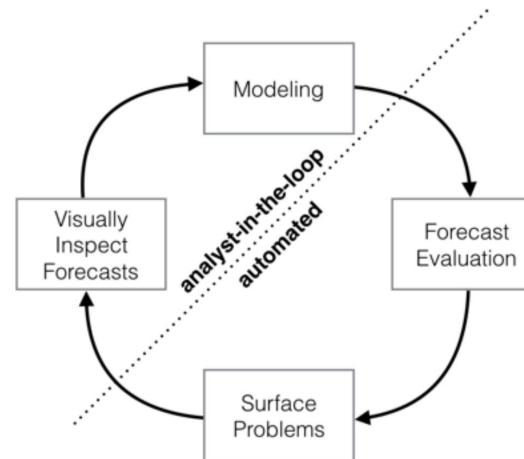
Sean J. Taylor^{*†}

Facebook, Menlo Park, California, United States
sjt@fb.com

and

Benjamin Letham[†]

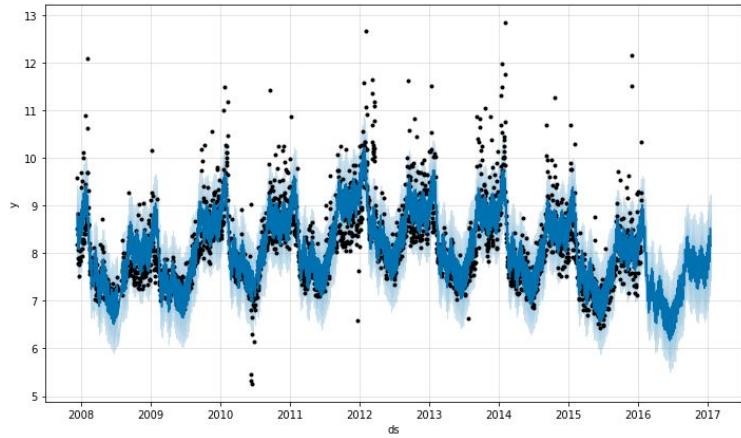
Facebook, Menlo Park, California, United States
bletham@fb.com



<https://peerj.com/preprints/3190/>

Prophet models interpretable time series components.

Introduction: Prophet



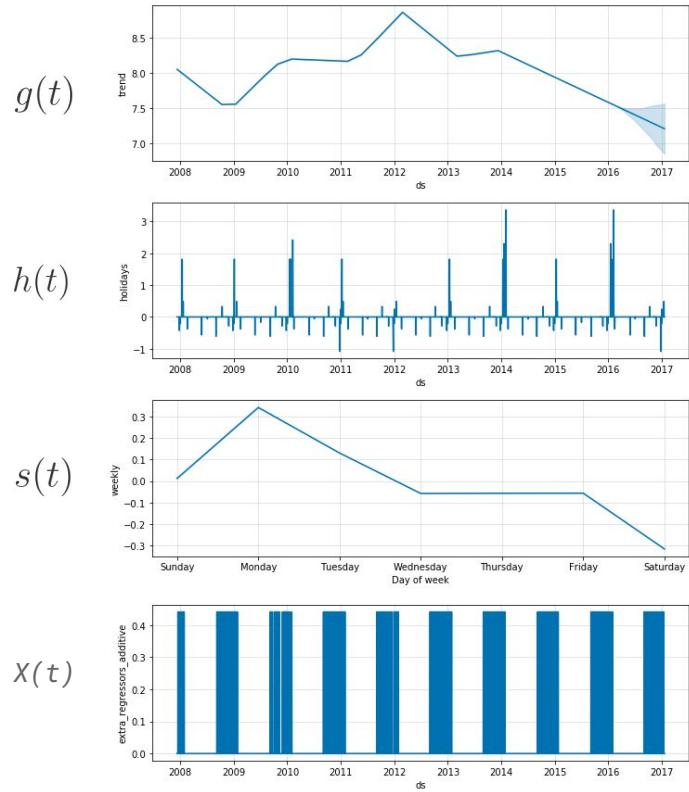
$$y(t) = g(t) + s(t) + h(t) + \epsilon_t.$$

$g(t)$ Trend

$h(t)$ Events

$s(t)$ Seasonality

$x(t)$ Regressors





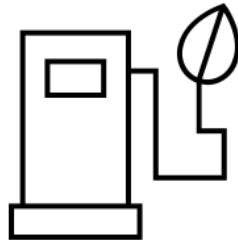
Extensible code.

Models suited for larger datasets.

Lagged input variables.

A car is more than an engine.

Introduction: NeuralPropet



Preprocessing



Diagnostics



Deployment

That's why we include helpful tools to streamline your workflow.

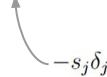
Current Model Components

S	<i>Seasonality</i>	-	<i>Sparsity / Regularization</i>
T	<i>Trend</i>	NN	<i>Nonlinear (deep) layers</i>
E/H	<i>Events / Holidays</i>	?	<i>Uncertainty</i>
AR	<i>Autoregression</i>		
Cov	<i>Covariates</i>		
X	<i>Regressors</i>		

Piecewise linear trend

- N changepoints
- Segmentwise independent
- Automatic changepoint detection
- Optional logistic growth

$$g(t) = (k + \mathbf{a}(t)^\top \boldsymbol{\delta})t + (m + \mathbf{a}(t)^\top \boldsymbol{\gamma}),$$



 $-s_j \delta_j$

Seasonality

- N Fourier terms
- Automatic yearly, weekly, daily

$$s(t) = \sum_{n=1}^N \left(a_n \cos\left(\frac{2\pi n t}{P}\right) + b_n \sin\left(\frac{2\pi n t}{P}\right) \right)$$

$$s(t) = X(t)\boldsymbol{\beta}.$$

$$X(t) = \left[\cos\left(\frac{2\pi(1)t}{365.25}\right), \dots, \sin\left(\frac{2\pi(10)t}{365.25}\right) \right]$$

Events / Holidays

- Automatic for given country
- Various user-defined formats

$$h(t) = Z(t)\kappa.$$

$$Z(t) = [\mathbf{1}(t \in D_1), \dots, \mathbf{1}(t \in D_L)]$$

(Future-Known) Regressors

- Single weight

(Lagged) Covariates

- AR-Net (y as target)

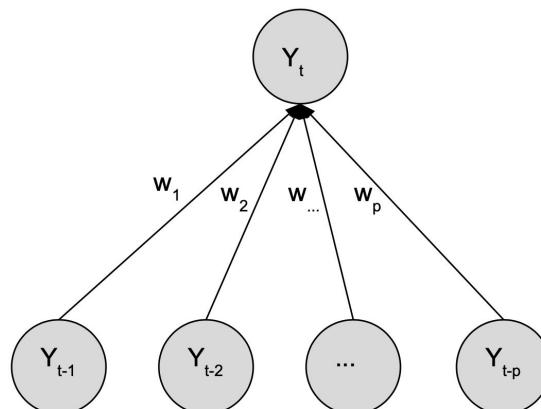
Auto-Regression

- AR-Net

AR-Net is a Neural Network for autoregressive time series.

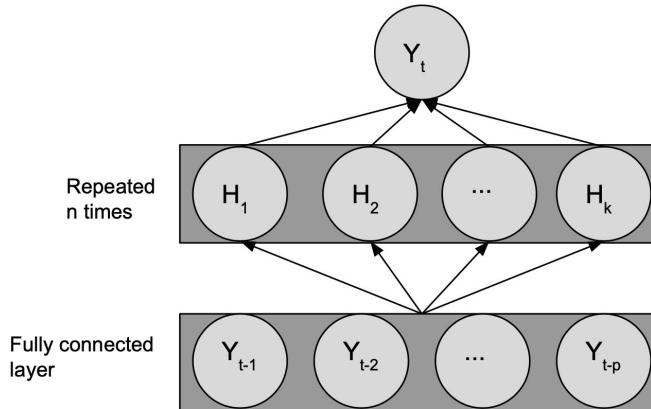
AR-Net(0)

Interpretable



AR-Net(n)

Stronger modeling ability



Skip estimating the AR process order.

Model: AR-Net

$$y_t = c + \sum_{i=1}^{i=p} w_i * y_{t-i} + e_t$$



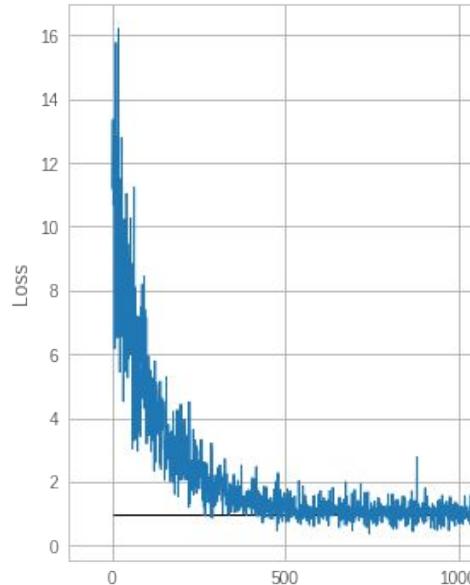
$$\min_{\theta} L(y, \hat{y}, \theta) + \lambda(s) \cdot R(\theta)$$

$$\lambda(s) = c_\lambda \cdot (s^{-1} - 1)$$

$$s = \frac{\hat{p}_{data}}{p_{model}}$$

$$c_\lambda \approx \frac{\sqrt{\hat{L}}}{100}$$

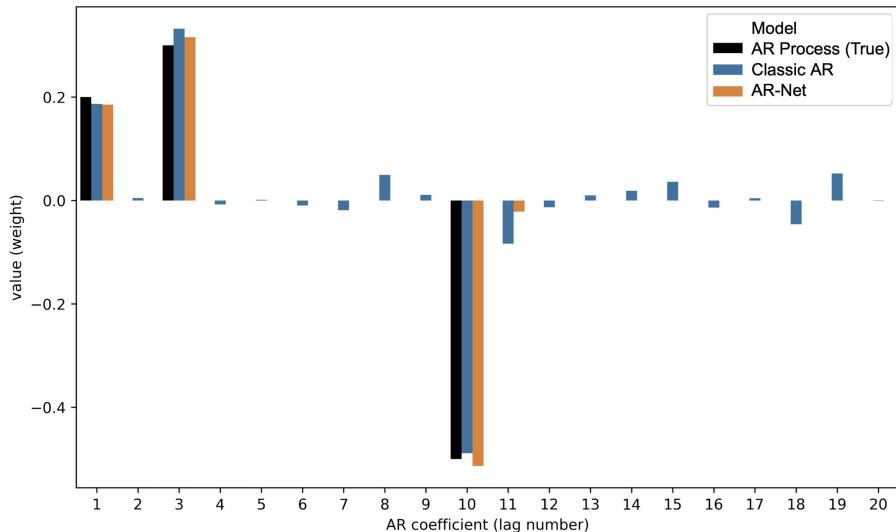
$$R(\theta) = \frac{1}{p} \sum_{i=1}^p \frac{2}{1 + \exp(-c_1 \cdot |\theta_i|^{\frac{1}{c_2}})} - 1$$



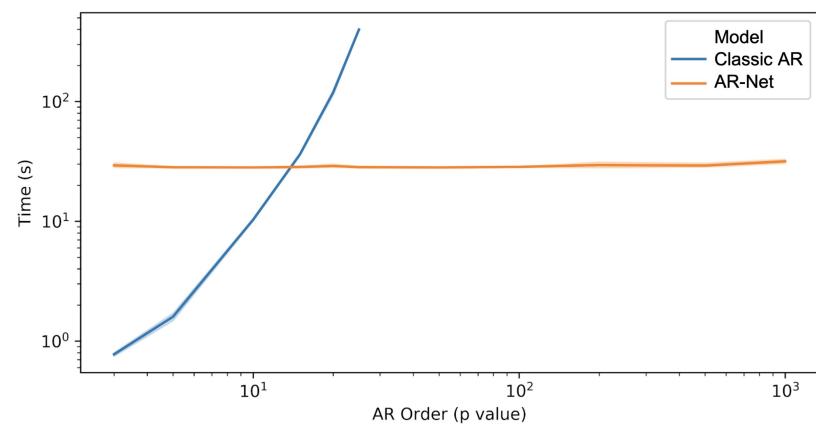
Trained with SGD (Adam)

Automatic AR-lag selection, yet faster.

Model: AR-Net



Automatic Sparsity



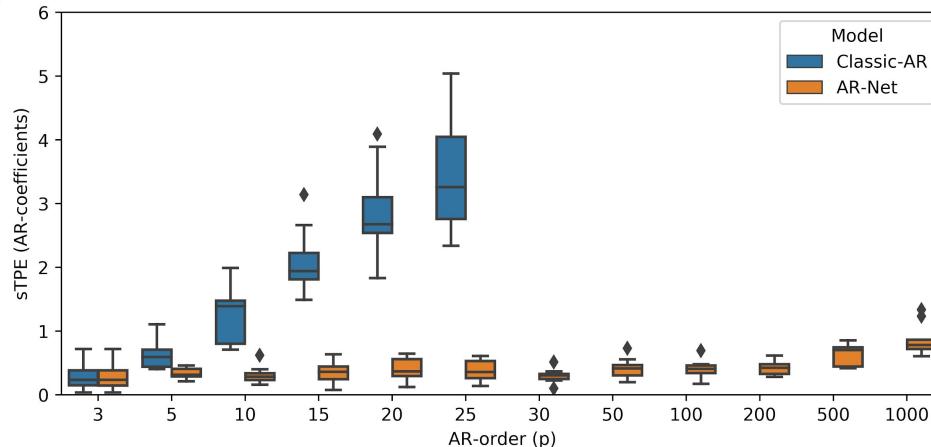
Quadratically faster

Sparse AR-Net surpasses Classic AR and scales to large orders.

Model: AR-Net

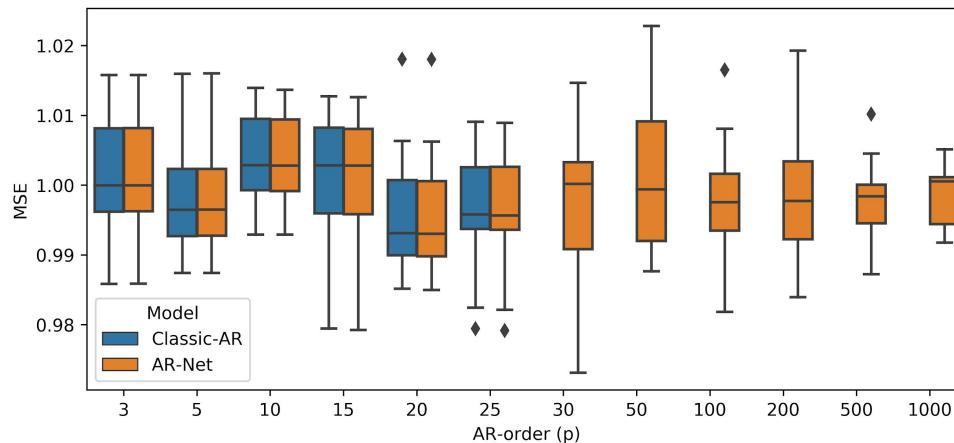
Closeness to true coefficients

$$sTPE = 100 \cdot \frac{\sum_{i=1}^{i=p} |\hat{w}_i - w_i|}{\sum_{i=1}^{i=p} |\hat{w}_i| + |w_i|}$$



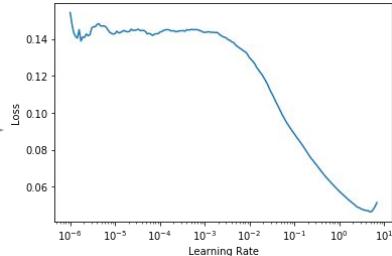
MSE loss on forecast target

$$MSE = \frac{1}{n} \sum_1^n (y_t - \hat{y}_t)^2$$

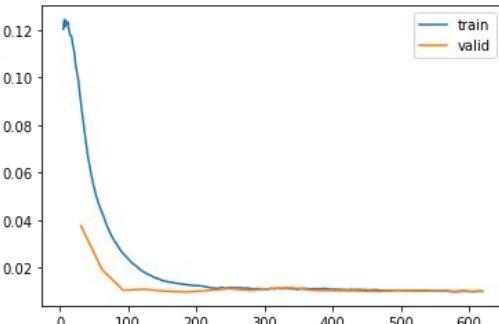


```
any time series  
learn = init_ar_learner(  
    series=df,  
    ar_order=3,  
    n_forecasts=1,  
    valid_p=0.25,  
    sparsity=1.0/2.0,  
)
```

```
lr = learn.lr_find()[0]  
learn.fit_one_cycle(n_epoch=10, lr)  
print(coeff_from_model(learn.model))
```



epoch	train_loss	valid_loss	mae	time
0	0.091714	0.037677	0.165344	00:00
1	0.043905	0.019270	0.112051	00:00
2	0.026220	0.010630	0.083301	00:00
3	0.018274	0.011065	0.084113	00:00
4	0.014546	0.010218	0.081142	00:01
5	0.013065	0.009898	0.080138	00:00
6	0.012065	0.010537	0.083343	00:01
7	0.011674	0.011297	0.085578	00:00
8	0.011395	0.010602	0.082547	00:00
9	0.011284	0.011390	0.085972	00:00
10	0.011243	0.011658	0.086935	00:00



building on open-source Fastai2, PyTorch, Python

```
ar coeff [[0.19132832, 0.3154733, -0.5822663, -0.049190696, 0.0042145597, 0.0015144324]]
```

Preprint on Arxiv

<https://arxiv.org/abs/1911.12436>

The screenshot shows the Arxiv preprint page for "AR-Net: A simple Auto-Regressive Neural Network for time-series". The page includes the title, authors (Oskar Triebe, Nikolay Laptev, Ram Rajagopal), submission date (27 Nov 2019), abstract, and several sections of the paper's content. On the right side, there is a sidebar with download options (PDF, Other formats), current browse context (cs.LG), change to browse by (cs, stat, stat.ML), references & citations, export citation, bookmark, and BibTeX links.

Code on Github

<https://github.com/ourownstory/AR-Net>

The screenshot shows the GitHub repository for "ourownstory / AR-Net". It displays the repository's main page with tabs for Code, Issues, Pull requests, Actions, Projects, Wiki, Security, Insights, and Settings. Below the tabs, it shows the master branch (1 branch, 1 tag), a list of commits, and a file tree. A callout points from the "v0.1" commit to the text "v 0.1: plain PyTorch". Another callout points from the "v1.0" commit to the text "v 1.0 built on Fastai2".

Commit	Message	Time
712bbe2	ourownstory update Readme	on Jun 11
v0.1	releasing v1.0 of AR-Net with simpler use, based on fastai2	7 months ago
v1.0	updated notebooks with cells run	6 months ago
AR-Net_paper.pdf	add paper to files	8 months ago
README.md	update Readme	4 months ago

Hands-on beta NeuralProphet

Basics

Trend

Seasonality

Events

Autoregression

Covariates / Regressors

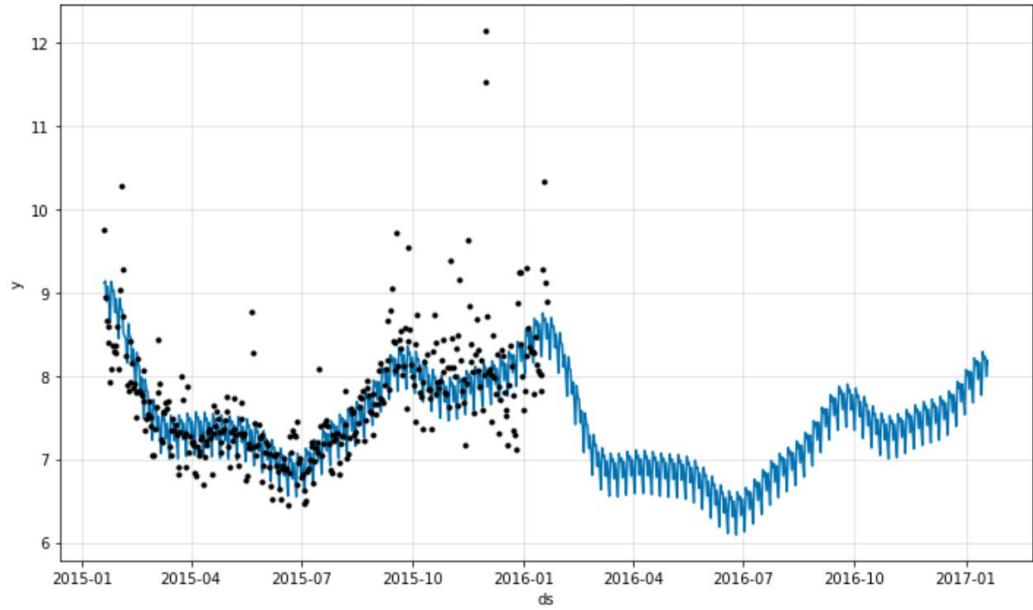
Few lines from data to forecast.

Hands-on

```
import pandas as pd
from neuralprophet.neural_prophet import NeuralProphet

df = pd.read_csv('../data/example_wp_log_peyton_manning.csv')

# linear time-dependent model
m = NeuralProphet()
metrics = m.fit(df)
future = m.make_future_dataframe(df, future_periods=365)
forecast = m.predict(future)
fig_fcst = m.plot(forecast[-730:])
```



```
| m = NeuralProphet()
```

Gentle learning curve.

NeuralProphet has smart defaults.

Only use what you need and understand.

```
n_forecasts=1,  
n_lags=0,  
n_changepoints=5,  
learning_rate=1.0,  
loss_func='Huber',  
normalize_y=True,  
num_hidden_layers=0,  
d_hidden=None,  
ar_sparsity=None,  
trend_smoothness=0,  
trend_threshold=False,  
yearly_seasonality='auto',  
weekly_seasonality='auto',  
daily_seasonality='auto',  
seasonality_mode='additive',  
seasonality_reg=None,  
data_freq='D',  
impute_missing=True,  
verbose=False,
```

Validate model performance without extra lines.

Hands-on

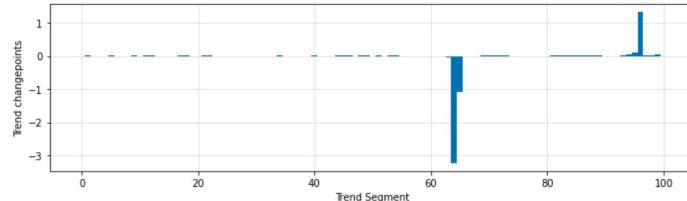
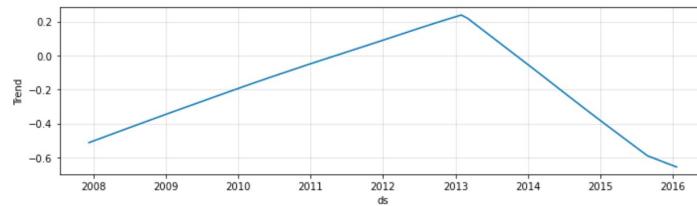
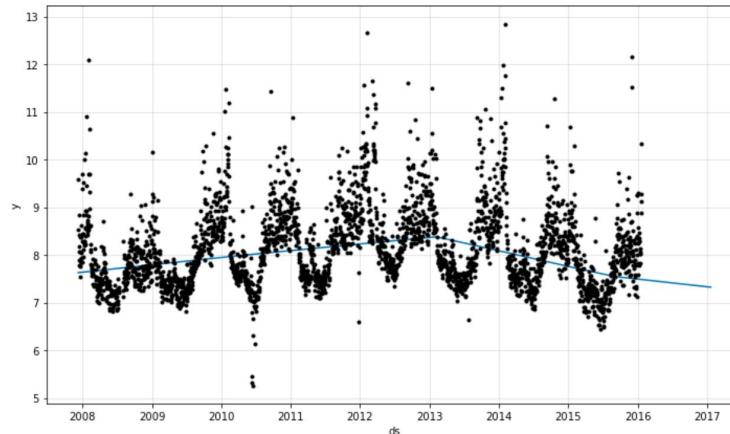
```
# or evaluate while training
m = NeuralProphet()
metrics = m.fit(df, validate_each_epoch=True, valid_p=0.2)
metrics.tail()
```

Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.

	SmoothL1Loss	MAE	RegLoss	SmoothL1Loss_val	MAE_val
35	0.163102	0.371323	0.0	0.485371	0.779465
36	0.161851	0.369609	0.0	0.368921	0.648736
37	0.161122	0.369219	0.0	0.366328	0.648230
38	0.168598	0.376638	0.0	0.348269	0.627886
39	0.167961	0.375777	0.0	0.362161	0.642699

```
# split manually
m = NeuralProphet()
df_train, df_val = m.split_df(df, valid_p=0.2)
train_metrics = m.fit(df_train)
val_metrics = m.test(df_val)
```

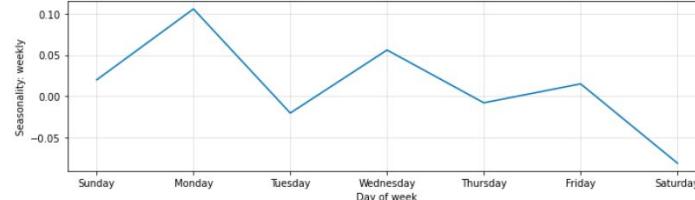
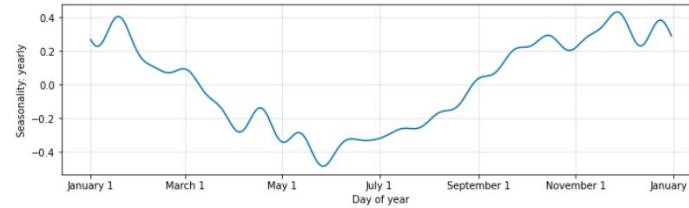
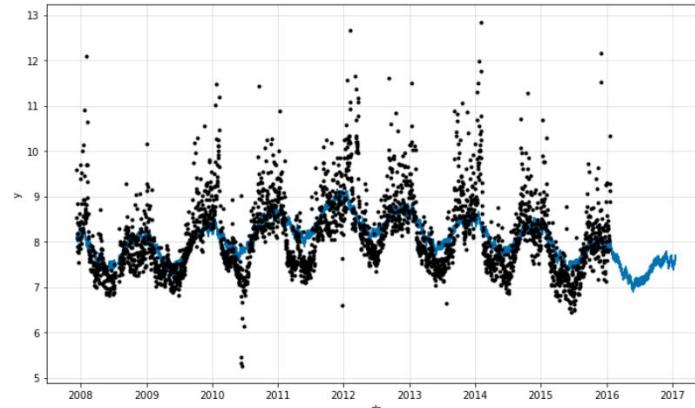
```
m = NeuralProphet(  
    n_changepoints=100,  
    trend_smoothness=2,  
    yearly_seasonality=False,  
    weekly_seasonality=False,  
    daily_seasonality=False,  
)  
metrics = m.fit(df)
```



Seasonality

Hands-on

```
m = NeuralProphet(  
    yearly_seasonality=16,  
    weekly_seasonality=8,  
    daily_seasonality=False,  
    seasonality_reg=1,  
)  
metrics = m.fit(df)
```



Events can be added in different forms.

Hands-on

```
superbowls = pd.DataFrame({'event': 'superbowl',
 'ds': pd.to_datetime(['2010-02-07', '2014-02-02', '2016-02-07'])})
```

```
events_df = pd.concat((superbowls, playoffs))
```

```
m = NeuralProphet()
```

```
m = m.add_country_holidays("US")
```

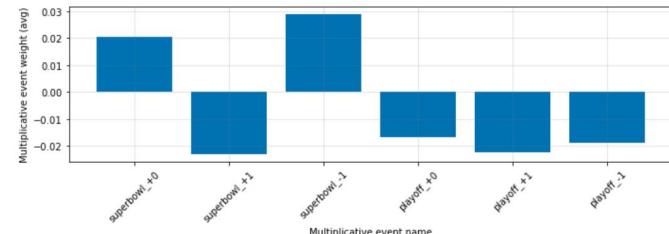
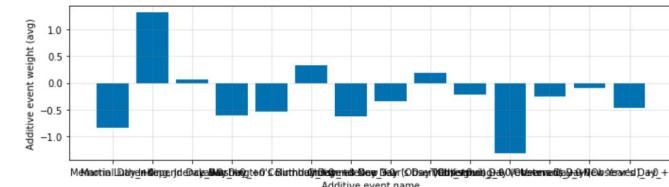
```
m = m.add_events(
    ["superbowl", "playoff"],
    lower_window=-1,
    upper_window=1,
    mode="multiplicative",
    regularization=0.5
)
```

```
history_df = m.create_df_with_events(df, events_df)
metrics = m.fit(history_df)
```

Disabling daily seasonality. Run NeuralProphet with daily_seasonality=True to override this.

```
future = m.make_future_dataframe(
    history_df,
    events_df,
    future_periods=30,
    n_historic_predictions=30
)
```

```
forecast = m.predict(future)
```



Auto-Regression

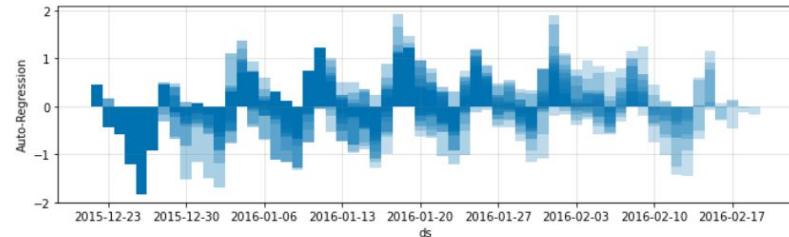
Hands-on

```
m = NeuralProphet(  
    n_forecasts=30,  
    n_lags=60,  
    ar_sparsity=0.1,  
    yearly_seasonality=False,  
    weekly_seasonality=False,  
    daily_seasonality=False,  
)
```

Autoregression,
here with sparsity

```
metrics = m.fit(df)  
future = m.make_future_dataframe(df, n_historic_predictions=30)  
forecast = m.predict(future)
```

```
fig_comp = m.plot_components(forecast[60:])
```



Lagged Covariates & Future Regressors

Hands-on

```
m = NeuralProphet(  
    n_forecasts=30,  
    n_lags=60,  
    ar_sparsity=0.1,  
    yearly_seasonality=False,  
    weekly_seasonality=False,  
    daily_seasonality=False,  
)
```

```
df = pd.read_csv('../data/example_wp_log_peyton_manning.csv')  
df['A'] = df['y'].rolling(7, min_periods=1).mean()  
df['B'] = df['y'].rolling(60, min_periods=1).mean()  
m = m.add_covariate(name='A')  
m = m.add_regressor(name='B')
```

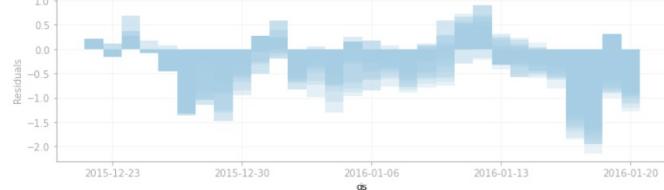
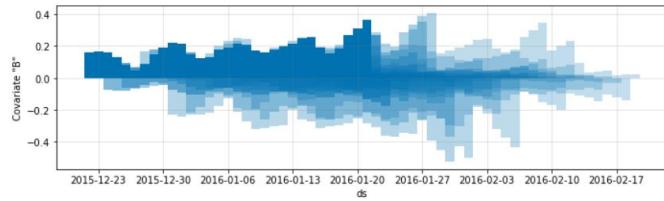
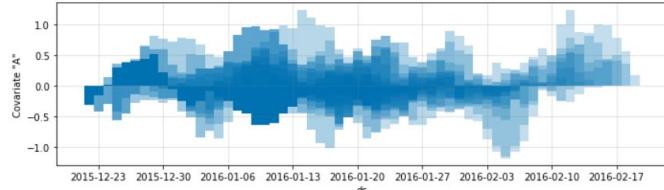
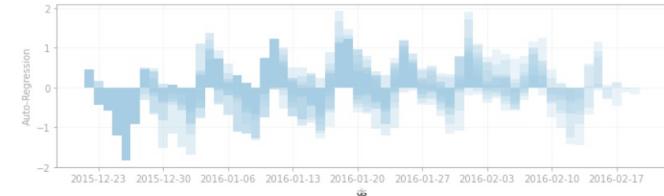
```
metrics = m.fit(df)  
future = m.make_future_dataframe(df, n_historic_predictions=30)  
forecast = m.predict(future)
```

Covariates and regressors are similarly added

Make it non-linear.

```
m = NeuralProphet(  
    n_forecasts=30,  
    n_lags=60,  
    learning_rate=1.0,  
    loss_func='Huber',  
    normalize_y=True,  
    num_hidden_layers=2,  
    d_hidden=64,  
)
```

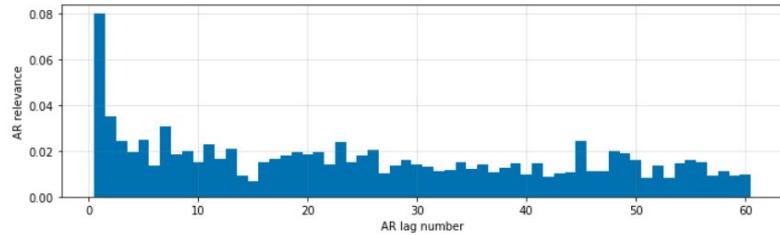
```
fig_comp = m.plot_components(forecast[60:])
```



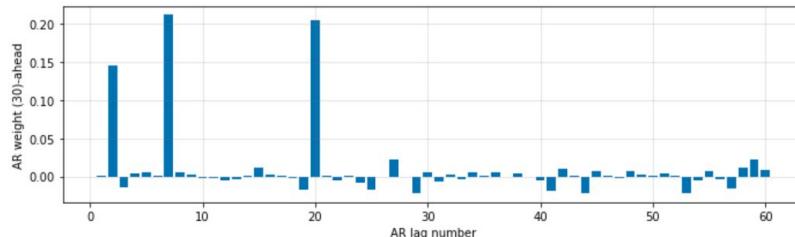
Focus on a specific forecast.

Hands-on

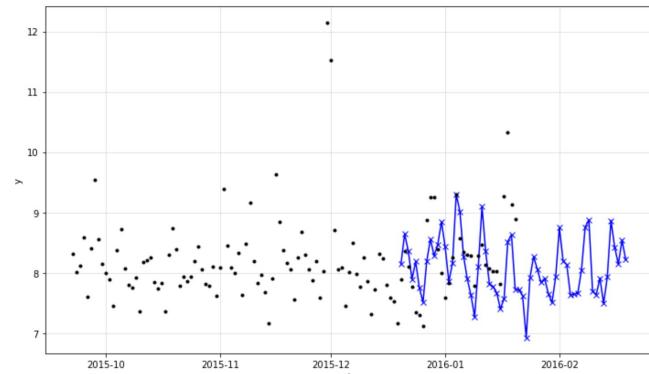
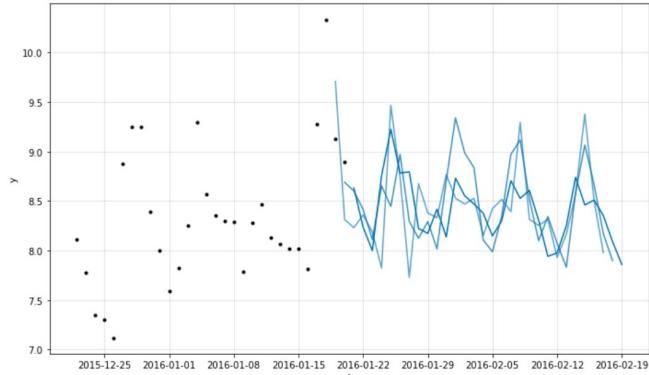
```
fig_param = m.plot_parameters()
```

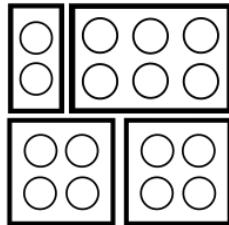


```
fig_param30 = m.highlight_nth_step_ahead_of_each_forecast(30).plot_parameters()
```



```
fig_prediction = m.plot_last_forecast(forecast[60:],  
include_previous_forecasts=2)
```





https://github.com/ourownstory/neural_prophet

Interpretable and customizable forecaster in PyTorch.

Extensible to future state-of-the-art in forecasting.

Workflow tools for diagnostics and more.

The future holds a lot of opportunities.



Upcoming:

- Beta release
- Documentation page
- Extensions (panel data, non-datetime data, ...)

Future:

V 1.0 Workflow utilities, robustify training, model extensions

V 2.+ Modularize, integrate other frameworks

V 3.+ Web GUI, port to other languages

THANK YOU, dear collaborators, sponsors and advisors!

Team



Oskar Triebel



Hansika Hewamalage



Nikolay Laptev



Ram Rajagopal



Riley De Haan



Gonzague Henri



Christoph Bergmeir



Italo Lima



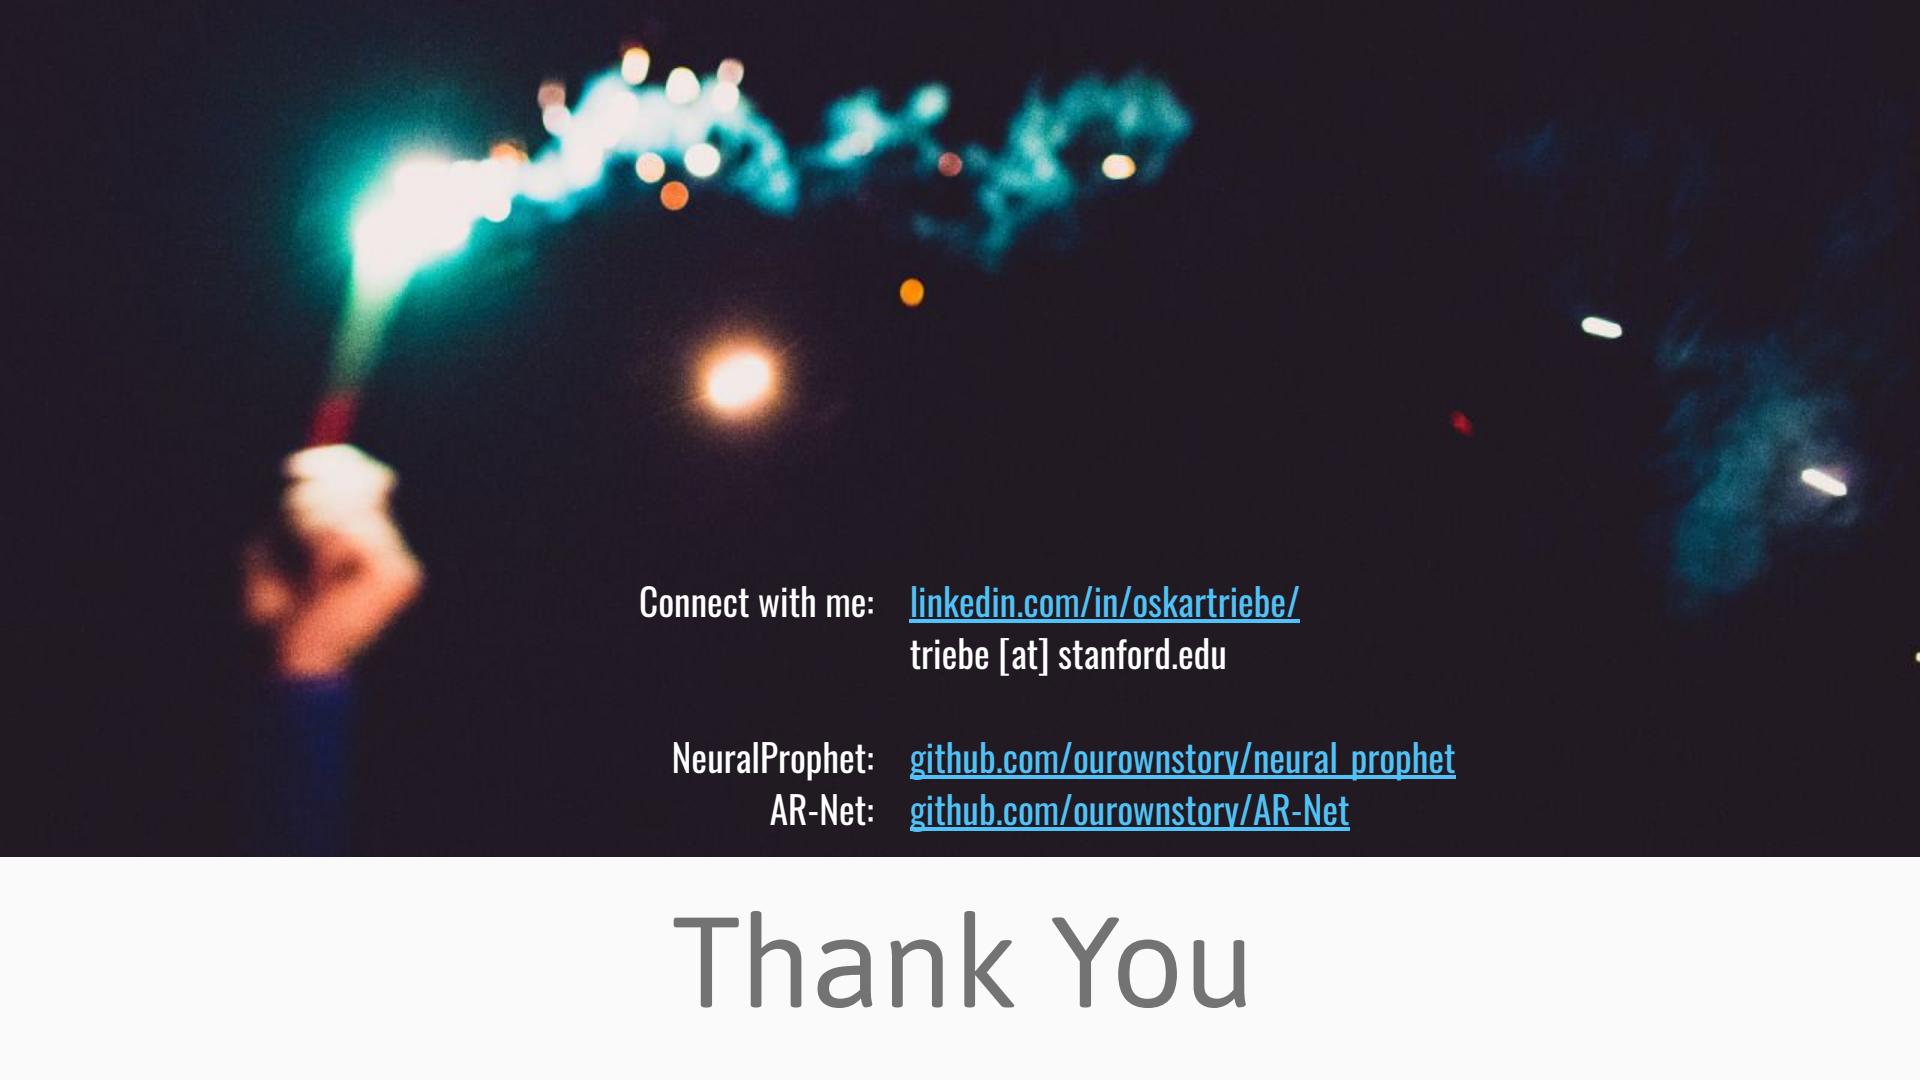
Caner Komurlu

Stanford University

TOTAL

MONASH
University

facebook



Connect with me: linkedin.com/in/oskartriebe/
triebe [at] stanford.edu

NeuralProphet: github.com/ourownstory/neural_prophet
AR-Net: github.com/ourownstory/AR-Net

Thank You